

Automated RCA Workflow — Reference Guide

Date: 2026-04-25

Audience: Managers, system engineers, plant data analysts

Source materials: `rca_metamodel.md` (causal categories A–L; **Pattern Recognition** questions), `orchestrators/rca_reasoning_orchestrator.py` (orchestrated run), `orchestrators/causality_engine_v32.py` (scoring and gates)

Show-and-tell tests: `src/dackar/RCA/tests/show_and_tell_test_plan.md`, `test_case_1` ... `test_case_7`

SE-focused summary: §1.1 — [Central elements](#). **Appendices:** **A** — [Pattern recognition](#); **B** — [LLMs in the pipeline](#); **C** — [Similar event search](#); **D** — [Show-and-tell test index](#).

This document is **descriptive** (what the pipeline is and what it does). It does not argue for or against particular licensing or deployment choices.

0. Document Conventions and Code-Documentation Sync

To keep this guide aligned with implementation over time, major sections are annotated with **HTML comments** (invisible in rendered Markdown) that a small script can parse:

- `<!-- @code: <path> | <class_or_fn> -->` — primary code anchor for the section
- `<!-- @schema: <path> -->` — JSON schema, when applicable
- `<!-- @status: implemented|partial|planned -->` — implementation state at last review
- `<!-- @reviewed: YYYY-MM-DD -->` — last time the section was verified against the repo

Example (as it appears in the source file, not the rendered view):

```
<!-- @code: orchestrators/rca_reasoning_orchestrator.py | RCAReasoningOrchestrator.run -->
```

`check_doc_code_sync.py` in this folder (`diagrams/april_25/`) validates `@code` / `@schema` anchors against files under `src/dackar/RCA/`, checks that Python **symbols** exist in the target module (class, method, or `def`), and is safe to use on CI. **Run (from DACKAR repo root):**

`python src/dackar/RCA/diagrams/april_25/check_doc_code_sync.py`, or from this directory: `python check_doc_code_sync.py` (optional: `--stale`, `--strict-warnings`). Bullets in the list above that use *placeholder* paths (not real `orchestrators/...` files) are **skipped** (not errors). Avoid putting additional `@code` HTML comments in the middle of a paragraph, or the checker will count them.

AP-913 completeness and regulatory mapping are **out of scope** for this revision; they will be added when the product owner locks the surface area.

1. Problem, Purpose, and What the Pipeline Produces

1.0 The problem

Large nuclear facilities generate an enormous volume of evidence during and after an equipment event: telemetry from the plant historian, sequence-of-events (SOE) records, alarm logs, condition reports, work orders, and maintenance history. Under schedule pressure, manual root-cause analysis is vulnerable to three well-documented failure modes:

- **Fixation** — early focus on the most salient signal, even when it is a consequence rather than a cause.
- **Incomplete coverage** — programmatic, common-cause, and human-performance contributors are under-weighted or missed entirely when time is short.
- **Inconsistent use of plant history** — relevant past events and operating experience exist in the plant record but are not systematically consulted.

The DACKAR RCA pipeline addresses these failure modes by constraining the hypothesis space to what plant data and the equipment model can support, scoring candidates in a deterministic and repeatable way, and surfacing data gaps and near-tie situations in structured artifacts the analyst and organization can review, challenge, and act on.

The analyst remains accountable for the final safety determination. The pipeline is an investigative and documentation engine, not a regulatory sign-off.

1.1 What the pipeline is designed to answer

A complete nuclear RCA must work through five investigation phases in order. The table below lists the key questions each phase must resolve and how the pipeline addresses them. Unanswered questions are surfaced as explicit gaps in the output — not silently dropped.

Investigation phase	Questions the pipeline must answer	How the pipeline addresses them
Scoping	What failed, on which equipment, under what operating conditions? Which safety functions were challenged or lost? What is the investigation boundary — systems, trains, time window?	The pipeline builds an initial scope from event data and available logs, tracks it as a versioned record, and flags out-of-boundary signals for analyst review (§5.1, §7).
Data and relations	What signals were anomalous before the event, and in what order? What maintenance, tests, or configuration changes occurred in the precursor window? Which anomalies are causes versus consequences? Are any relevant data sources missing, degraded, or of insufficient resolution?	Signal ordering is established by placing telemetry anomalies, alarms, and SOE events on a shared timeline. Every data source — present, missing, or degraded — is recorded in the run manifest with a sensitivity flag indicating whether its absence could change the ranking (§5.3, §5.4, §6).
Pattern recognition	Is there a degradation trend preceding the event, and over what timescale? Is this a first occurrence or a recurrence — and if recurrence, did prior corrective actions work? Is there evidence of common-cause failure across multiple trains or components? Does the anomaly signature match a known failure mode?	Temporal patterns are scored against plant history and known failure-mode signatures. Past events from the equipment model are matched to active candidates. Common-cause failure is a scored dimension when vendor and lot records are available (§2.4, Appendix A).
Hypothesis generation and ranking	What are the plausible root causes, ranked by evidence strength? Which causal categories have been ruled out, and why? Are there near-tie hypotheses the evidence cannot discriminate between? Is the top-ranked hypothesis physically consistent with the operating conditions at the time?	Candidates are generated from the equipment model (not free text), scored on five weighted dimensions, filtered through three hard gates (physical plausibility, timeline consistency, barrier logic), and then checked against a dual composite/evidence-floor threshold. Ruled-out candidates carry documented reasons. Past-event profile matches are returned as a separate event analogs pool. Near-ties and sensitivity to missing data are flagged explicitly (§5.6, §5.8).
Conclusion	What corrective actions address the proximate, contributing, and root cause levels? What barriers failed, and why? Was human performance a cause or contributor? What gaps remain unresolved? What should be monitored to verify corrective action effectiveness?	The RCA card organizes findings at three causal depths (proximate → contributing → root), includes a barrier assessment and human performance block when evidence supports them, and lists unresolved gaps with the evidence that would close them (§5.12, §6).

Coverage enforcement: For each event, the pipeline must either generate and score at least one candidate in each of the 12 causal categories (A–L, §1.5) or explicitly document in the RCA card why a category was ruled out. Missing coverage is an attention flag, not a silent omission.

1.2 What the pipeline produces

Every run produces a **result bundle** under a unique run identifier. The two artifacts most relevant to managers and system engineers are:

rca_card — the primary analyst-facing document. Contains: ranked hypotheses with evidence citations; corrective action recommendations organized by causal depth (proximate, contributing, root); barrier assessment; human performance findings when applicable; and unresolved gaps with the specific evidence that would close each one.

run_manifest — the audit and review record. Contains: which data sources were present, missing, or degraded; how sensitive the current ranking is to each missing source (sensitivity table); what scope decisions were made and when; and what review actions the analyst must complete before the record can be closed.

Supporting artifacts (`causality_candidates` , `evidence_bundle` , `tskr_patterns` , `barrier_analysis` , `ishikawa_matrix` , and others) provide the evidence trail behind the card. They are described by audience in **§8**; the complete input and output list is in **§2**.

1.3 What the pipeline does not replace

- Physical inspection, laboratory analysis, or on-site equipment examination
- NRC event reporting determinations or AP-913 regulatory sign-off
- The analyst’s final safety and quality determination — the pipeline provides structured inputs to that determination, not a substitute for it
- Fleet or industry operating-experience databases — these require a site-configured adapter (see **Appendix C**)
- A trained neural model — all candidate scoring is rule-based and deterministic; weights and thresholds are transparent and configurable (§5.6)

1.4 How the pipeline is built: main components

The pipeline is assembled from a set of purpose-built components. Each one is responsible for a distinct part of the investigation, so any finding in the RCA card can be traced back to the component — and from there to the data — that produced it. The table maps each component to the §1.1 investigation phases it serves.

Component	What it does	Investigation phase served
Equipment and failure mode model	Holds a structured representation of plant equipment, component connections, failure modes, and safety function definitions. Every hypothesis is anchored to a specific component and failure mode already in this model — not invented from raw signals alone.	Scoping; Hypothesis generation and ranking
Data coverage assessor	Checks which data sources were supplied, records each as present, missing, or degraded, and flags whether any missing source could change the ranking.	Data and relations
Temporal sequencer	Places telemetry anomalies, alarms, and SOE events on a shared timeline to separate precursors (potential causes) from responses (potential consequences), and scores degradation trends per failure mode.	Data and relations; Pattern recognition
Hypothesis ranking engine	Generates candidates from the equipment model, scores each on five dimensions (structural fit, temporal support, telemetry match, documentary evidence, maintenance program posture), and applies three binary elimination gates. Eliminated candidates are retained in a documented ruled-out list.	Pattern recognition; Hypothesis generation and ranking
Document evidence retriever	Searches condition reports, work orders, procedures, FMEAs, and operating experience documents for text relevant to each active candidate. Retrieved passages adjust candidate scores and are cited in the RCA card.	Hypothesis generation and ranking; Conclusion
Operating experience matcher	Scores plant history for past events on the same or similar equipment, and optionally queries fleet or industry databases when a site adapter is configured.	Pattern recognition; Conclusion
RCA card synthesizer	Assembles the final structured report from ranked candidates, evidence passages, and barrier status, organized at three causal depths: proximate, contributing, and root cause.	Conclusion

Every intermediate result is saved to the run record, and the **run manifest** (§1.2) tracks which components ran, what data each had access to, and what analyst review actions remain open.

1.5 Causal categories (A–L)

Every hypothesis the pipeline generates is assigned to one of twelve causal categories. Together they define the coverage requirement for a complete nuclear RCA: for each event the pipeline must either produce at least one scored candidate in each applicable category, or explicitly document in the RCA card why that category does not apply. Authoritative sub-item definitions are in `diagrams/april_25/rca_metamodel.md`.

Cat	Name	What it captures	Causal depth	Primary driver data
A	Equipment-internal	Intrinsic material, mechanical, electrical, or instrumentation degradation of the component itself	Proximate cause	Telemetry anomalies, FMEA failure mode nodes in KG
B	Required support unavailable	Failure of an ancillary system — power, cooling, lubrication, instrument air, control signal — that the primary component depends on	Proximate cause	Ancillary system tags, KG support-system edges
C	Upstream influence / common cause	Inlet conditions outside design basis; lot or batch common-mode defects across multiple trains	Proximate cause	Process topology, <code>vendor_supply_chain_records</code> , KG CCF delta
D	Downstream influence	Backpressure, blocked discharge, or unstable demand imposed on the component outlet by downstream systems	Proximate cause	Process topology, telemetry downstream tags
E	Operating context / mission demand	Equipment operated outside its design envelope: overload, off-design transients, prolonged standby, or excessive cycling	Proximate cause	<code>operational_context</code> (mode, train configuration, demand history)
F	External hazards	Conditions from outside the plant process boundary: thermal environment, flooding, seismic, fire, EMI	Proximate cause	<code>environmental_monitoring</code> , OE documents, evidence bundle
G	Human execution	Operator or maintenance action that deviated from a correct procedure or configuration baseline	Contributing cause	Work order text in Chroma, <code>training_records</code> (qualification data)
H	Design deficiency	Equipment performs as designed, but the design itself is inadequate for actual service conditions	Contributing cause	Design documents and engineering evaluations (ECA, ECR) in Chroma
I	Configuration and change control	Work was executed correctly against a wrong, outdated, or unrevised configuration baseline	Contributing cause	<code>configuration_change_records</code> , procedure documents in Chroma
J	Inspection / testing program gap	Surveillance program not designed to detect this failure mode at its actual progression rate	Contributing cause	PM compliance records, <code>pm_compliance</code> artifact
K	Vendor and supply chain	Specification was correct but the delivered item did not meet it; includes lot-level and batch defects	Contributing cause	<code>vendor_supply_chain_records</code> , BULLETIN / MANUAL documents in Chroma
L	Systemic / organizational weakness	Programmatic or cultural root cause that allowed contributing causes to persist across events	Root cause	OE documents, prior CR effectiveness reviews, <code>similar_event_list</code> recurrence data

The causal depth column maps to the AP-913 framework. Corrective actions that address only the proximate cause without also addressing contributing and root cause levels will not satisfy regulatory expectations.

1.6 Candidates vs. causal categories

Causal categories (A–L) are classes of causal mechanisms — they define the *type* of cause and ensure investigation coverage. Coverage enforcement operates at the category level: did the pipeline generate at least one candidate in each applicable category, or explicitly rule it out?

A **candidate** is a specific, instance-level hypothesis: a particular failure mode on a particular component, at a particular point in the event causal chain, attributed to a primary causal category. Ranking and evidence assessment operate at the candidate level.

Internally, each candidate is represented as a **4-tuple**:

Dimension	Values	What it means
Component	Component ID from the equipment model	Which physical item failed or contributed
Failure mode	Failure mode ID from the equipment model (FMEA-derived)	The specific mechanism by which that component failed
Causal category	A–L (§1.5)	The class of cause — determines which data sources and gates apply
Chain position	initiating / contributing / consequence	Where in the causal chain this candidate sits

`chain_position` distinguishes a candidate that directly caused the event (`initiating`) from one that enabled or worsened it (`contributing`) or that is a downstream effect of it (`consequence`). The timeline consistency gate eliminates consequence candidates — a signal that follows the event cannot be its cause.

Each candidate additionally carries a **composite score** (weighted sum of five sub-scores), the results of the three hard gates, a score rationale, and, after the evidence refinement step, citations into the document corpus. The composite score is what drives ranking; the 4-tuple is what makes a candidate traceable to a specific physical item in the plant.

One category can generate multiple candidates (e.g. several different failure modes on the same pump, or the same failure mode on different pumps across trains). One candidate belongs to exactly one primary category.

2. Workflow at a Glance

2.1 Input data inventory

Two inputs are always required: **event** (event identity, asset ID, time window, symptom text) and **telemetry_summary** (pre-processed anomaly records per signal — pattern type, severity, time window). Every other input is optional; when absent it is recorded in the run manifest as `not_assessed` and the causal categories it would have informed are flagged as data-limited. The pipeline always completes; it never silently drops a coverage gap.

Note on telemetry: the pipeline reasons over anomaly records — it does not perform signal processing on raw time-series. Anomaly detection quality upstream directly limits what the pipeline can conclude.

Note on telemetry unavailability: `telemetry_summary` is a hard required input; the pipeline cannot accept `None` in its place. If pre-processed anomaly records are not available, the alarm log provides a partial substitute for **interval ordering only** — alarm activations are placed on the Allen relation map (§2.3) and can contribute to timeline classification. However, alarm records carry no anomaly `pattern_type` or `severity`, so **TSKR scoring is entirely absent** when `telemetry_summary` is missing. All candidates will have a zero temporal sub-score, pattern recognition (§2.4 Questions 1 and 4) will not run, and the RCA card will reflect heavily degraded rankings. The run manifest will flag all TSKR-dependent causal categories as `data-limited`.

Note on SOE and protection logic: these two sources are tightly coupled. If SOE is present but protection logic is absent, the timeline and barrier gates operate in degraded mode and the run manifest escalates an `analyst_decisions_required` flag.

Inputs are grouped below by data family, following the same structure as the causal taxonomy (§1.5).

Event identity (required)

Data element	What it provides to the pipeline	Causal categories	Schema
Event	The investigation anchor. Carries: <code>event_id</code> , <code>asset_id</code> , <code>timestamp_start</code> , <code>timestamp_end</code> , <code>event_type</code> , <code>actuation_type</code> , and symptom description. Every other input is cross-validated	A–L	<code>event.json</code>

Data element	What it provides to the pipeline	Causal categories	Schema
	against the <code>event_id</code> and <code>asset_id</code> fields here — a mismatch aborts the run. Required.		

Pre-built pipeline artifacts (optional, checkpoint-resume)

These optional inputs allow the pipeline to skip re-computing one or more steps by injecting outputs from a prior run. When supplied they bypass the corresponding computation stage; when absent the stage runs normally. They are not plant data sources — they are previously computed pipeline state.

Artifact	Bypasses	Schema
<code>signal_evidence</code>	Step 2a — signal propagation chain build	<code>signal_evidence.json</code>
<code>tskr_patterns</code>	Step 2b — TSKR temporal scoring	<code>tskr_patterns.json</code>
<code>causality_candidates</code>	Steps 3–4 — candidate generation and initial ranking	<code>causality_candidates.json</code>
<code>evidence_bundle</code>	Step 5 — Chroma evidence retrieval	<code>evidence_bundle.json</code>

Structured and model data

Data element	What it provides to the pipeline	Causal categories	Schema
Equipment model	Components, topology, failure modes, safety function definitions, and plant past events — the typed hypothesis space. Every candidate is anchored to a node in this model. Built from Neo4j at run time or supplied as a pre-built snapshot.	A–L	<code>kg_context.json</code>
Protection logic	Trip setpoints, permissive logic diagrams, and barrier / safety function states at event time. Required for the barrier elimination gate and for interpreting SOE records.	A, B, F	<code>protection_logic_context.json</code>
Configuration change records	Engineering change notices, setpoint change packages, and work order history — change control traceability in the precursor window.	H, I	<code>configuration_change_records.json</code>

Time-series and event data

Data element	What it provides to the pipeline	Causal categories	Schema
Telemetry summary	Pre-processed anomaly records per sensor: pattern type (e.g. gradual drift, step rise), severity, and anomaly time window. Required.	A–F	<code>telemetry_summary.json</code>
Sequence of events (SOE) log	Millisecond-resolution discrete event records from the SOE recorder. Feeds the timeline consistency gate and the temporal sequencer. Pair with protection logic for full barrier analysis.	A, B, F	<code>soe_log.json</code>
Alarm log	Alarm activations with timestamps and priorities. Extends temporal coverage when SOE is absent or sparse.	A–F	<code>alarm_log.json</code>
Environmental monitoring	Ambient temperature, humidity, seismic indicators, and grid disturbances at event time.	F	<code>environmental_monitoring.json</code>

Maintenance and operations data

Data element	What it provides to the pipeline	Causal categories	Schema
Operational context	Operating mode, power level, recent operator actions, and shift narrative at event time. Feeds the operating-point scoring dimension (Category E modifier). Operator shift logs should be included here — shift turnover notes, operator observations, and equipment status entries from the shift log are expected as part of <code>shift_narrative</code> and <code>recent_alarms</code> . There is no separate <code>operator_shift_log</code> input; if shift logs become a first-class data source they will require a dedicated schema.	E, G	<code>operational_context.json</code>
PM compliance	PM schedule adherence, overdue intervals, and as-found / as-left inspection results. Auto-built from CMMS export rows when not supplied directly.	A, J	<code>pm_compliance.json</code>
Condition reports and work orders (CR / WO)	Event-scoped corrective maintenance records from the plant CMMS (e.g. Maximo). Not a direct <code>run()</code> input — fetched internally when a CMMS adapter is configured on the orchestrator. CR/WO records contribute to: (1) the initial scope <code>component_ids</code> at Step 0; (2) the plant past-event pool (up to 12 records merged into <code>kg_context.past_events</code> at Step 1); (3) the KG document neighborhood. When no CMMS adapter is configured this data is absent and flagged in <code>run_manifest.artifacts.cmms_context</code> .	A, G, J, L	<code>cmms_context.json</code> (internal artifact)

Document and institutional knowledge


Data element	What it provides to the pipeline	Causal categories	Schema
Document corpus	Pre-indexed condition reports, work orders, procedures, FMEAs, ECAs, and industry OE documents held in the vector store. Retrieved per candidate to raise or lower evidence scores and populate RCA card citations. Supplied as a pre-built bundle or queried live from Chroma.	All	<code>evidence_bundle.json</code> (pre-built) or live Chroma retrieval (§3.2)

Supply chain and vendor data

Data element	What it provides to the pipeline	Causal categories	Schema
Vendor and supply chain records	Lot numbers, vendor certifications, and receipt inspection records. Feeds the Category C common-cause structural score when a batch or lot defect is suspected.	C, K	<code>vendor_supply_chain_records.json</code>
Training records	Personnel qualification status and training recency. Provides evidence for human performance and programmatic findings; does not by itself establish root cause.	G, L	<code>training_records.json</code>

2.2 Workflow steps (system engineer map)

The pipeline executes as a fixed, deterministic sequence of seven steps aligned with the investigation phase structure of the causal taxonomy (§1.5). Each step has a defined scope, consumes specific inputs, runs a deterministic mechanism, and produces artifacts that downstream steps build on. The table below is the SE reference map; §5 provides full per-step detail.



Simplified logical data flow. *event* is a universal input: its *asset_id* scopes the KG neighborhood, its *timestamp_start/end* is the reference interval for all Allen and TSKR scoring, and its *event_type* and *actuation_type* are query terms for the similar event search. Optional inputs (SOE, alarms, PM, vendor records, protection logic, etc.) are not shown.

Step	Scope	Key inputs	Core mechanism	Key c
0 – Initialize	Prove inputs are structurally valid and freeze run identity before any analysis begins.	event (required), telemetry_summary (required), operational_context , alarm_log , soe_log , protection_logic_context , configuration_change_records	JSON schema validation of all supplied inputs; cross-checks (asset ID consistency, telemetry coverage); initial investigation scope assembled from event data and logs (scope version 0, no boundary filter active); pm_compliance auto-built from CMMS export rows if not supplied.	run_context (sc guard flags), inpr pm_compliance
1 – Equipment model and data context	Materialize the typed hypothesis space and enrich it with CMMS records, past event history, and per-signal anomaly chains.	event , telemetry_summary , run_context , operational_context , pm_compliance	KG queried (or snapshot loaded): components, topology, failure modes, safety functions, past events. KG governance assessed; red state can abort per policy. CMMS adapter (if configured) augments with live CR/WO records. Past events enriched with temporal metadata. Signal-to-component tag mapping builds scored anomaly propagation chains.	kg_context , kg. signal_evidence (if adapter configu
2 – Temporal analysis and OE search	Establish the full temporal picture and operating-experience context before hypothesis ranking: search equipment history, score per-failure-mode temporal support, build the event-wide signal ordering map, and identify similar events. Four sub-steps (2a–2d):	event , telemetry_summary , kg_context (with past events from KG and CMMS), signal_evidence , alarm_log , soe_log , operational_context	2a – Equipment history temporal search: For each component/asset in the KG scope, past CRs, WOs, surveillance records, and operator logs (loaded into <code>kg_context.past_events</code> at Step 1) are filtered to those within the precursor window and tagged by window tier. This builds the recurrence pool used by TSKR in 2b. 2b – TSKR per-failure-mode scoring: For each failure mode, computes onset lag, duration profile, recurrence count, and novelty flag from anomaly windows and the precursor-window record pool.	tskr_patterns (support, recurrence flags), allen_rel wide signal orderi similar_event_l industry matches)

Step	Scope	Key inputs	Core mechanism	Key c
			Novelty flag set when <code>recurrence_count == 0</code> and no prior signal match exists. 2c — Allen relation map: All anomaly intervals, alarm activations, and SOE point events are classified relative to the triggering event using Allen's interval algebra (OVERLAPS, CONTAINS, PRECEDES, DURING, FOLLOWS) — separating causal precursors from consequence followers for the whole event. 2d — Similar event identification: Plant past events scored on five dimensions (component match, failure mode, event type, actuation type, precursor window overlap); top matches returned. Fleet and industry databases queried via <code>SimilarEventAdapter</code> when configured (§2.6).	
3 / 3.5 — Pattern recognition	Identify recurrence and novelty across two evidence registers: the documentary record (past CRs, WOs, and confirmed RCA findings) and the signal record (anomaly shapes, alarm sequences, SOE transitions). Two sub-steps:	<code>tskr_patterns</code> , <code>kg_context</code> (past events enriched with CRs/WOs/RCA findings from Step 1), <code>signal_evidence</code> , <code>alarm_log</code> , <code>soe_log</code>	Step 3 — Documentary pattern recognition: Examines <code>kg_context.past_events</code> records — structured prior events carrying <code>root_cause_label</code> , <code>finding_status</code> , and <code>causal_explanation</code> extracted from CRs, WOs, and confirmed RCA reports (injected via CMMS at Step 1 and KG historical data). For each prior event, checks component match, failure mode match, and resolution status. <code>_build_historical_support_channels</code> produces counts: same-component, same-failure-mode, unresolved, CMMS-injected. TSKR uses this pool to compute <code>recurrence_count</code> and <code>recurrence_trend</code> per failure mode. <i>Note: raw-text retrieval from CRs, WOs, RCAs, and ECAs in Chroma is Step 5 — Step 3 operates on structured metadata only.</i> Step 3.5 — Signal pattern recognition: <code>_build_signal_lessons_learned</code> classifies each TSKR pattern as matched (historical support found: <code>recurrence_count > 0</code> or <code>history_score ≥ threshold</code>) or novel (<code>novel_pattern = True</code> , no prior signal match). Input sources: telemetry anomaly windows, alarm log activations, SOE point events.	Recurrence and n <code>tskr_patterns</code> ; <code>signal_lessons_</code> vs novel pattern c assembled into ru 6)
4 — Candidate generation	Generate all plausible failure-mode hypotheses, apply initial scoring and hard elimination	<code>event</code> , <code>kg_context</code> , <code>tskr_patterns</code> (<i>carries Steps 2b + 3 outputs</i>),	Candidates generated per failure mode from the equipment model — never from free text. Five-dimension weighted	<code>causality_candi</code> refine; contains FI event analogs, an

Step	Scope	Key inputs	Core mechanism	Key c
and initial ranking	gates, enforce approved scope boundary. Directly depends on Steps 2b and 3 via <code>tskr_patterns</code> : <code>causality_engine.generate()</code> immediately indexes <code>tskr_patterns</code> on entry and uses it to compute the temporal sub-score for every candidate. Because <code>tskr_patterns</code> carries the Step 3 recurrence outputs (<code>recurrence_count</code> , <code>recurrence_trend</code> , <code>novel_pattern</code> , <code>unresolved_recurrence_count</code>), Step 4 implicitly consumes both the TSKR temporal scoring (Step 2b) and the documentary recurrence analysis (Step 3) through a single artifact. <code>signal_evidence</code> (Step 2a propagation chains) is not passed to <code>generate()</code> — it enters at Step 5 refinement only.	<code>operational_context</code> , <code>pm_compliance</code> , <code>run_context</code>	scoring: structural (KG topology, CCF delta), temporal (TSKR recurrence + Allen blend from <code>tskr_index</code>), telemetry (anomaly fit), evidence (document availability), governance (PM posture). Three binary hard gates applied in order: physical plausibility, timeline consistency, barrier logic. A dual-threshold check then retains only candidates that meet both the composite score minimum and the pre-evidence floor. Eliminated candidates retained in a documented ruled-out list with reason codes. Past-event profile matches returned as a separate event analogs pool. Scope boundary filter applied if an approved scope revision is active.	
5 — Evidence retrieval and refinement	Retrieve documentary evidence per candidate, re-score the ranked list, and optionally expand context if coverage gaps are detected.	<code>causality_candidates</code> (v1), <code>event</code> , <code>kg_context</code> , <code>allen_relation_map</code> , <code>protection_logic_context</code> , <code>vendor_supply_chain_records</code> , <code>training_records</code> , <code>environmental_monitoring</code>	Chroma vector store queried per active candidate → scored passages (supporting, contradicting, contextual). Second scoring pass (<code>refine_with_evidence</code>) blends evidence scores, Allen relation signals, and protection logic into composite scores and re-runs hard gates. Optional auto-reentry re-runs Steps 1–5 with expanded context when coverage gaps are detected and <code>enable_auto_reentry</code> is set.	<code>evidence_bundle</code> <code>causality_candi</code> refine), <code>causality_candi</code> <i>(for diff and reentry)</i>
6 — Conclusion and finalization	Assemble the final RCA card and barrier assessment; validate outputs; produce the run manifest.	All prior artifacts (including <code>similar_event_list</code> from Step 2d); <code>protection_logic_context</code> , <code>vendor_supply_chain_records</code> , <code>training_records</code> , <code>environmental_monitoring</code>	Optional Ishikawa 6M bucketing of evidence themes. Barrier analysis aggregates safety-function status from protection logic and hard gate results. RCA card assembled from ranked candidates, evidence citations, similar event findings, and barrier results at three causal depths (proximate → contributing → root). Attention flags applied (rank inversions, governance state, data quality). Output schema validation. Scope expansion signals detected from Allen map and TSKR state and injected into <code>run_context</code> for the next run. Run manifest assembled with data coverage, sensitivity table, scope state, and review hooks.	<code>rca_card</code> , <code>barr</code> . <code>ishikawa_matrix</code> <code>output_validation</code> <code>run_manifest</code> , <code>reentry_execution</code>

Scope revision: Scope is a versioned record initialized at Step 0 (version 0, open — no candidates filtered). If the analyst accepts an expansion suggestion from a prior run, scope version ≥ 1 is active on re-run and Step 4 applies the approved boundary filter before evidence retrieval. See [§2.5.3](#) for the scope revision workflow.

2.3 Allen temporal relations

The Allen relation map (pipeline Step 2, §2.2) classifies every observed signal, alarm, and discrete event relative to the triggering event using a subset of Allen's interval algebra. The result — `allen_relation_map` — is a single shared artifact built once and reused by the hypothesis ranking engine (Step 4 temporal blend), the evidence refinement pass (Step 5), the scope expansion signal detector (Step 6), and the run manifest.

Relations and causal interpretation

The pipeline implements five of the thirteen Allen relations, selected for their RCA relevance. Each relation carries a base score that serves as a prior in temporal scoring; these scores are further refined downstream by latency alignment and severity weighting.

Relation	A relative to triggering event B	Base score	Causal candidate
OVERLAPS	A started before B and was still active at B onset — degradation present at the moment of failure	0.90	Yes
CONTAINS	A started before B and ended after B — a long-running latent condition that encompasses the event	0.85	Yes
PRECEDES	A ended before B started — classic causal lead-time; the anomaly was already resolved when the event occurred	0.75	Yes
DURING	A started at or after B onset — anomaly appeared inside or after the event; likely a consequence, not a cause	0.30	No
FOLLOWS	A starts after B ends — temporal contradiction; cannot be a cause	0.10	No

Relations are evaluated in the order `FOLLOWS` → `PRECEDES` → `CONTAINS` → `OVERLAPS` → `DURING`. An `epsilon_hours` tolerance (default **0.5 h**) absorbs timestamp noise and near-simultaneous boundary cases — boundaries within epsilon are treated as touching. An `interval_type` parameter controls endpoint interpretation (closed / open / half-open) for edge cases.

`PRECEDES` , `OVERLAPS` , and `CONTAINS` set `causal_candidate = True` on the output node. `DURING` and `FOLLOWS` do not.

Input sources and interval representation

Both intervals and instantaneous events (time points) are supported. Point events are modeled as a degenerate closed interval with `start = end` . The table below describes the three input sources and how each is mapped to an interval.

Source	Input fields consumed	Interval representation	Always a point event?	Clock sync guar
Telemetry anomalies	<code>telemetry_summary.signals[].anomaly_window.start</code> / <code>.end</code> (or <code>anomaly_start</code> / <code>anomaly_end</code>)	<code>[anomaly_start, anomaly_end]</code> ; if no end field, treated as point (<code>end = start</code>)	No — duration present when historian provides a window	None; anomaly timestamps assumed reliable
Alarm log entries	<code>alarm_log.alarms[].activated_at</code> (start); <code>cleared_at</code> or <code>acknowledged_at</code> (end)	<code>[activated_at, cleared_at]</code> ; if not yet cleared, treated as point (<code>end = start</code>)	No — most alarms have a clear time	Yes — if <code>alarm_log.quality.clock</code> is <code>False</code> , relation forced to "unknown" with score 0.0

Source	Input fields consumed	Interval representation	Always a point event?	Clock sync guard
SOE records	soe_log.records[].timestamp	Always a point event (start = end = timestamp): SOE records represent instantaneous state transitions	Always	Yes — if soe_log.quality.clock_sync == False , relation forced to "unknown" with score 0.0; logs capped at max_soe_nrecords (default 200)

Output per node

Each processed event produces one node in allen_relation_map.nodes[] carrying:

Field	Description
node_type	anomaly , alarm , or soe_record
component_id	Component linked to the signal or record (when available)
interval_start / interval_end	ISO timestamps; interval_end is null for point events
is_point_event	true for SOE records and for anomalies/alarms with no duration
allen_relation_to_event	One of the five relations above, or "unknown" on clock failure
allen_base_score	Base relevance score (0.10–0.90), 0.0 when unknown
causal_candidate	true when relation is PRECEDES, OVERLAPS, or CONTAINS

The map-level summary includes timeline_consistent (true when no FOLLOWS nodes exist), counts of causal_nodes and contradiction_nodes , and earliest_causal_onset — these feed attention flags on the RCA card and the run manifest.

How Allen scores enter the scoring pipeline

During Step 5 (refine_with_evidence), allen_base_score is blended into each candidate's temporal sub-score:

temporal_refined = 0.75 × TSKR_score + 0.25 × allen_base_score

The blend is one-directional — it can only raise, not lower, a candidate's temporal score. A FOLLOWS relation on a component matched to an active candidate sets temporal_contradiction = True , which the timeline consistency gate uses to eliminate that candidate from the ranked list.

2.4 Pattern recognition

Pattern recognition in this pipeline is not a single component or stage — it is a cross-cutting capability distributed across Steps 2 through 6. The four questions below, drawn from the causal taxonomy metamodel, are each answered by a different combination of input data, pipeline steps, and output fields. The SE's role is to verify that the right input data is present for each question and to interpret the corresponding output artifacts correctly.

Question 1 — Is there a degradation trend, and over what timescale?

What the pipeline looks for: whether a signal anomaly was building up gradually before the event, how long it persisted, and how it was positioned relative to the event onset.

	Detail
Input data	telemetry_summary.signals[].anomaly_window (start/end timestamps, anomaly type); alarm_log ; soe_log
How it works	Each anomaly type in telemetry_summary — e.g. gradual_drift , sustained_exceedance , step_rise , oscillation — characterizes the shape of the degradation. The TSKR scorer (Step 2) computes onset lag hours and a duration profile per failure

	Detail
	mode. The Allen relation (§2.3) classifies whether the anomaly was active at event onset (<code>OVERLAPS</code>), long-running and encompassing the event (<code>CONTAINS</code>), or already resolved before the event (<code>PRECEDES</code>).
What it produces	<code>tskr_patterns.patterns[].onset_lag_hours</code> — how far before the event the anomaly began; <code>tskr_patterns.patterns[].duration_profile</code> — characterization of how long the pattern persisted; <code>allen_relation_map.nodes[].allen_relation_to_event</code> — the temporal position of each signal. <code>OVERLAPS</code> and <code>CONTAINS</code> nodes are the strongest indicators of a pre-existing degradation trend.
Pipeline steps	Step 2 (TSKR, Allen map); Step 4 (temporal sub-score); Step 5 (Allen blend into refined score)

Question 2 — Is this a first occurrence or a recurrence? Did prior corrective actions work?

What the pipeline looks for: whether the current event matches past events on the same or similar equipment, and whether the pattern is genuinely novel (no prior signal match) or a recurrence of a known failure sequence.

	Detail
Input data	<code>kg_context.past_events</code> (enriched with temporal metadata at Step 1); <code>vendor_supply_chain_records</code> (for lot-level recurrence); optionally fleet/industry OE via a configured adapter (§2.6)
How it works	At Step 1, past events in the equipment model are tagged with <code>in_precursor_window</code> and window tier. At Step 3 / 3.5, each TSKR pattern is checked for historical matches: if <code>recurrence_count == 0</code> and no prior signal match exists, the pattern is flagged <code>novel_pattern = True</code> . At Step 6, the similar event list scores each plant past event on five dimensions (component match, failure mode match, event type, actuation type, precursor window overlap) and ranks the top matches.
What it produces	<code>tskr_patterns.patterns[].novel_pattern</code> (True = no prior occurrence found); <code>tskr_patterns.summary.has_novel_patterns</code> ; <code>signal_lessons_learned</code> in <code>run_manifest.artifacts</code> — separates matched patterns (prior causal explanation known) from novel ones; <code>similar_event_list</code> with confidence-weighted plant past events and <code>any_plant_match</code> summary flag.
Pipeline steps	Step 1 (past event temporal enrichment); Step 3/3.5 (novelty flags); Step 6 (similar event list, signal lessons learned)

Question 3 — Is there evidence of common-cause failure across trains or components?

What the pipeline looks for: whether the same failure mode is appearing on multiple components simultaneously, pointing to a shared upstream cause (Category C) or a vendor batch defect (Category K) rather than an independent per-component failure.

	Detail
Input data	<code>kg_context</code> (topology showing which components share support paths or are from the same lot/vendor); <code>vendor_supply_chain_records</code> (lot numbers, batch certifications, receipt inspection records)
How it works	The equipment model captures which failure modes carry a common-cause indicator. When a candidate belongs to Category C (upstream influence shared across trains), the hypothesis ranking engine adds a CCF structural delta to the structural sub-score, proportional to the <code>common_cause_score</code> computed from the KG topology. When vendor records are supplied, lot overlap across affected components raises this score further.
What it produces	<code>causality_candidates[].primary_causal_category == "C"</code> candidates in the ranked list; <code>causality_candidates[].scores.ccf_score</code> and <code>.ccf_note</code> on every candidate (0.0 for non-CCF candidates); <code>causality_candidates[].common_cause</code> block with structural CCF reasoning. A Category C candidate ranked highly is the pipeline's signal that CCF warrants investigation.
Pipeline steps	Step 1 (KG topology, vendor records loaded); Step 4 (CCF structural delta in initial scoring); Step 5 (CCF re-evaluated in refinement pass)

Question 4 — Does the anomaly signature match a known failure mode?

What the pipeline looks for: whether the combination of anomaly type, severity, affected component, and temporal position is consistent with a specific failure mode already in the equipment's FMEA.

	Detail
Input data	telemetry_summary (anomaly type, severity); kg_context (failure mode definitions, component-FM bindings from FMEA); tskr_patterns (per-FM temporal support and confidence)
How it works	The hypothesis ranking engine generates candidates by binding each FMEA failure mode to its parent component — only failure modes already in the equipment model are considered. The structural sub-score measures how well the component topology and failure mode type fit the observed event. The telemetry sub-score measures how well the anomaly severity and type match the failure mode signature. The temporal sub-score from TSKR measures how well the observed anomaly timing matches the expected onset lag and duration profile for that failure mode. The three sub-scores combine into a composite that ranks signature strength.
What it produces	causality_candidates ranked list — the top candidates represent the strongest signature matches to known failure modes; causality_candidates[].scores.structural_score , .telemetry_score , .temporal_score and their score_rationale entries explain which dimensions drove the ranking. A high composite score with high temporal and telemetry sub-scores indicates a strong signature match.
Pipeline steps	Step 2 (TSKR per-FM confidence and support); Step 4 (structural + telemetry + temporal sub-scores in initial ranking); Step 5 (evidence sub-score adds documentary support to the signature match)

Where pattern recognition results surface for the analyst

Pattern recognition outputs are not consolidated into a single artifact. The analyst sees them across:

- **tskr_patterns** — degradation trend and signature match (Questions 1 and 4)
- **allen_relation_map** in `run_manifest` — signal ordering and timeline consistency (Question 1)
- **signal_lessons_learned** in `run_manifest.artifacts` — novelty vs. recurrence summary (Question 2)
- **similar_event_list** — past event matches ranked by similarity (Question 2)
- **causality_candidates** — CCF scores and Category C candidates (Question 3); composite scores and score rationale for all four questions
- **rca_card.attention_flags** — novel pattern flag, near-tie flag, and sensitivity flags surface pattern recognition gaps that require analyst attention

Appendix A covers limitations and improvement directions for the pattern recognition mechanisms described in §2.4.

2.5 Analyst interaction points

The pipeline is designed to run to completion without pausing mid-execution. Analyst interaction happens at two well-defined points: **after each run** (reading the outputs and deciding what to do next) and **between runs** (updating `run_context` with decisions that take effect on the next execution). There is no interactive mid-run prompt.

The four interaction patterns are described below.

2.5.1 Reading attention flags

After every run the `rca_card.executive_summary.analyst_attention_flags` list collects all conditions the pipeline could not resolve on its own. The analyst reads these flags as part of reviewing the card. The pipeline populates them from seven independent checks:

Flag trigger	What it means for the analyst
Rank inversion — pre-evidence top candidate differs from post-evidence leader	Evidence quality or KG coverage may be distorting the ranking; verify both before accepting the card
KG governance warning — equipment graph health is yellow or red	Hypothesis space may be incomplete or stale; check KG update status before writeback

Flag trigger	What it means for the analyst
High CR match-failure rate — recurrence pool has many unmatched condition records	Recurrence ranking is likely understated; check whether relevant past events are missing from the KG
Signal propagation cascade truncated — feedback loop detected in Stage B.5	Causal chain topology has concurrent-cause loops; review the chain manually before signing off
Out-of-boundary anomaly signals — anomalies exist outside the current investigation boundary	Signals were excluded by scope; consider whether they represent upstream causes (see §2.5.3)
Metamodel category coverage unknown — one or more A-L categories could not be assessed	Pipeline cannot confirm coverage for those categories; analyst must determine applicability (high-impact: B, F, I, L)
Ishikawa not performed — <code>ishikawa_matrix</code> was absent	Human performance and organizational factor branches were not evaluated; assess manually if the event warrants it

2.5.2 Reading review hooks and deciding the next step

After every run the `run_manifest.review_hooks` block tells the analyst what the pipeline thinks should happen next. The key fields:

Field	Values	Meaning
<code>next_step</code>	<code>"writeback"</code>	All gates passed; run is ready to write to plant records
<code>next_step</code>	<code>"analyst_review"</code>	Outputs are valid but human sign-off is required before writeback
<code>next_step</code>	<code>"validation_remediation"</code>	A hard gate or governance check failed; remediation required before proceeding
<code>requires_human_review</code>	<code>true</code>	At least one condition (uncited claim, evidence gate failure, decision pending, severity floor missed) requires human review
<code>writeback_ready</code>	<code>true</code>	All schema, citation, evidence, severity, and decision checks passed and no degraded-run conditions exist
<code>analyst_decisions_required</code>	List of strings	Specific items the analyst must resolve (e.g. pending scope expansions, SOE/PLC pairing violation) before the run can progress
<code>coverage_acknowledgement_required</code>	<code>true</code>	Data coverage is partial or missing; analyst must set <code>analyst_review.coverage_degraded_acknowledged = true</code> in the card before writeback is allowed
<code>hard_abort_required</code>	<code>true</code>	Strict red-state governance is active and the KG is in a red state; the run must not proceed until governance remediation is complete

If `writeback_ready` is `false` the analyst reads `degraded_reasons` (a plain-text list) to understand exactly which conditions are blocking writeback.

2.5.3 Scope revision — accepting, deferring, or rejecting expansion suggestions

When Step 6h detects signals that fall outside the current investigation boundary it writes **expansion suggestions** into `run_context.scope_management.expansion_suggestions`. Each suggestion carries a `signal_id`, a list of `suggested_component_ids`, and an initial `analyst_decision` of `"pending"`.

The analyst resolves each suggestion before the next run using `resolve_expansion_suggestion`:

```
updated_context = orchestrator.resolve_expansion_suggestion(  
    run_id=run_id,  
    run_context=run_context,  
    signal_id="<signal_id from suggestion>",  
    decision="accepted",          # or "deferred" or "rejected"  
    rationale="Pump upstream of failed valve - in scope",  
)
```

Decisions and their effects:

Decision	Scope version	Effect on next run
"accepted"	Incremented by 1	apply_scope_revision is called immediately; added component IDs are merged into the approved boundary; Step 4 (§5.6) will apply the boundary filter on re-run
"deferred"	Unchanged	Suggestion stays pending; pipeline re-runs with the same boundary
"rejected"	Unchanged	Suggestion is marked rejected; pipeline re-runs with the same boundary

The versioned scope history is preserved in `scope_management.scope_revisions[]` . Version 0 is always the initial open boundary (no filter). Version ≥ 1 means at least one accepted revision is active.

Multi-run pattern (TC-7): Run 1 → analyst reviews expansion suggestions → `resolve_expansion_suggestion` → Run 2 with narrowed or expanded boundary → analyst compares ranked candidates before and after scope change.

2.5.4 Data backfill — supplying missing inputs and re-running

When the sensitivity table in `run_manifest.artifacts.data_coverage_summary` flags a source as "partial" or "missing" , it means that source, if supplied, could materially change the candidate ranking or gate outcomes. The analyst's options:

Situation	Action
Missing source is available	Supply the JSON input on the next run; the pipeline will re-score with the additional data
SOE present but protection_logic_context absent	This triggers a paired-data "violated" flag and adds an item to <code>analyst_decisions_required</code> ; supply <code>protection_logic_context</code> or explicitly accept degraded barrier-gate operation in <code>analyst_review</code>
Missing source is genuinely unavailable	Note it in <code>analyst_review</code> ; the <code>unresolved_gaps</code> field in the card is populated for exactly this purpose

Coverage acknowledgement: if data coverage is degraded and the analyst has confirmed the run should proceed anyway, set `analyst_review.coverage_degraded_acknowledged = true` in the card JSON before the next run. This clears `coverage_acknowledgement_required` from the review hooks.

2.6 Fleet, plant, and industry OE (similar events)

The pipeline searches for similar past events across three tiers — plant, fleet, and industry — during Step 2d (`_build_similar_event_list`), before the RCA card is synthesized. Results feed into `rca_card.unresolved_gaps` and the card narrative to contextualize whether the current event is a known recurrence or a novel occurrence.

The **plant tier always runs** (in-memory, no external call). The **fleet and industry tiers run only when a `SimilarEventAdapter` is injected** via `orchestrator.set_similar_event_adapter(adapter)` . Without an adapter, `similar_event_list.status` is reported as "partial" .

2.6.1 Query construction

Before querying any tier the orchestrator builds a shared set of query terms from the top-3 causality candidates (configurable via `step2d_query_top_n_candidates`):

Query term	Source
asset_id	Event input

Query term	Source
component_ids	component_id fields of top-N candidates
failure_mode_ids	failure_mode_id or canonical_tuple.failure_mode of top-N candidates
event_type	Event input
actuation_type	Event input

These terms are stored verbatim in `similar_event_list.query_terms` for audit traceability.

2.6.2 Plant tier

The plant tier is always active and requires no external connection. It scores the `past_events` array from `kg_context` (the equipment model) against the current event. Four plant data sources contribute, each at a different stage:

Plant data source	Contents	Role in the plant tier
Equipment KG (<code>kg_context.past_events</code>)	Structured equipment-level event history loaded from the Neo4j graph neighborhood	Primary pool scored directly by <code>_query_plant_past_events</code> ; always present when the KG was built with historical data
CMMS (<code>cmms_context.cr_records</code> , <code>cmms_context.wo_records</code>)	Condition reports and work orders from the plant Computerized Maintenance Management System	Up to 12 records merged into <code>kg_context.past_events</code> before scoring (configurable via <code>cmms_past_event_injection_max</code>); also contributes component IDs to the initial scope boundary
Alarm log (<code>alarm_log.alarms</code>)	Timestamped alarm transitions with <code>system</code> field	Does not enter <code>past_events</code> directly; at Step 0 each alarm's <code>system</code> field populates <code>scope_snapshot.system_boundary</code> , shaping which KG neighborhood — and therefore which past events — are loaded for the run
SOE log (<code>soe_log.records</code>)	Sequence-of-events records with <code>component_id</code> field	Does not enter <code>past_events</code> directly; at Step 0 each record's <code>component_id</code> is added to <code>scope_snapshot.component_ids</code> , which together with alarm-derived system boundaries defines the initial KG query scope

The net effect is that alarm and SOE records shape the investigation boundary at Step 0, which determines the KG neighborhood, which determines what is in `kg_context.past_events` for plant-tier matching. Alarm and SOE records also appear in `signal_lessons_learned.summary.input_sources` as provenance for the pattern scoring step.

Scoring is deterministic and multi-dimensional:

Dimension	Weight	Match condition
Component match	0.40	<code>matched_component_ids</code> overlap with current candidates
Failure mode match	0.25	<code>matched_failure_mode_ids</code> overlap with top-5 candidate FMs
Event type match	0.15	<code>event_type</code> equals current event
Actuation type match	0.10	<code>actuation_type</code> equals current event
Precursor window boost	0.10	<code>in_precursor_window = true</code> on the past event

The confidence weight is `min(1.0, raw_score × 1.00)` (plant tier multiplier = 1.0). The top-5 results by confidence weight are returned (configurable via `step2d_plant_top_n`). Each returned record carries `source_db = "plant_kg"` .

2.6.3 Fleet and industry tiers

Fleet and industry queries run via the `SimilarEventAdapter` protocol. The concrete production implementation is `LLMOEAdapter` (`adapters/llm_oe_adapter.py`), which calls a fine-tuned LLM REST API.

Tier	What is queried	Data sources backing the LLM	Confidence multiplier	source_db label
Fleet	Utility fleet operating experience records	Utility-peer event databases (owner-group operating experience)	0.80	"fleet_oe"
Industry	Broad industry OE database	INPO SOER, EPRI technical reports, NRC LERs	0.60	"inpo_epri_nrc"

The orchestrator calls `adapter.query(level="fleet", ...)` and `adapter.query(level="industry", ...)` separately, each capped at 5 results and 10-second timeout. If either call fails, that tier is added to `similar_event_list.summary.degraded_tiers` and the status is set to `"partial"` .

Configuring the adapter:

```
from adapters.llm_oe_adapter import LLMOEAdapter

adapter = LLMOEAdapter(
    fleet_url="https://oe-api.example.com/fleet",
    industry_url="https://oe-api.example.com/industry",
    api_key=os.environ["OE_API_KEY"],
    model_name="oe-finetuned-v1", # fine-tuned on INPO/EPRI/NRC corpus
)
orchestrator.set_similar_event_adapter(adapter)
```

The confidence multiplier discounts apply to fleet (0.80) and industry (0.60) records *after* the adapter returns its raw scores, reflecting increasing distance from plant-specific context. Plant records are not discounted (multiplier 1.0).

2.6.4 Output artifact

All three tiers are merged into `similar_event_list` :

```
similar_event_list
├─ status           "complete" | "partial"
├─ query_terms      audit record of what was searched
├─ summary
│   ├─ plant_count
│   ├─ fleet_count
│   ├─ industry_count
│   ├─ total_count
│   ├─ any_plant_match
│   └─ degraded_tiers  list of tiers that failed or had no adapter
├─ events[]         merged list, each with source_level, confidence_weight,
│                   source_db, failure_signature, root_cause_label,
│                   resolution, lessons_learned_ref, contributing_categories
└─ provenance       adapter class name, generator
```

`status = "complete"` only when an adapter was present **and** both fleet and industry tiers completed without errors. `any_plant_match = true` indicates at least one plant-tier event was found matching the current event dimensions — the strongest signal for recurrence assessment.

2.7 Pipeline outputs

Every run produces a result bundle keyed by a unique `run_id` . Outputs fall into three categories: the **primary analyst-facing document** (`rca_card`), the **audit and review record** (`run_manifest`), and a set of **intermediate artifacts** that form the evidence trail behind the card.

All artifacts are persisted to the artifact store (default: `FileArtifactStore`) as JSON files under `<run_id>/<artifact_name>.json`. The two primary outputs (`rca_card`, `run_manifest`) are always written; intermediate artifacts are written as each pipeline step completes.

2.7.1 `rca_card` — primary analyst-facing document

The `rca_card` is produced by `RuleValidatedRCASynthesizerV31` (Step 6d / §5.12) and is the document the analyst reads, challenges, and signs off on before writeback. Its top-level structure:

```
rca_card
├── executive_summary
│   ├── decision_status      "candidate_ready" | "review_required" | "insufficient_evidence"
│   ├── primary_conclusion    plain-text summary sentence
│   ├── confidence_label      "high" | "medium" | "low" | "speculative"
│   └── analyst_attention_flags[] conditions requiring analyst review (see §2.5.1)
├── primary_hypothesis
│   ├── candidate_id, cause_label, fm_id, hypothesis_type
│   ├── narrative            explanation of why this is the primary cause
│   ├── why_primary[]        reasons this ranks above alternatives
│   ├── uncertainties[]      what remains unconfirmed
│   ├── composite_score (float)
│   ├── confidence_label
│   └── citations[]          evidence IDs supporting this hypothesis
├── alternatives[]           ranked list; each has candidate_id, cause_label,
                             composite_score, reason_not_primary, citations[]
├── contributing_causes[]    each with contribution_type
                             ("contributing" | "enabling" | "escalating"),
                             rationale, citations[]
├── evidence[]              all evidence items used; each has evidence_id,
                             source_type, support_role
                             ("supporting" | "contextual" | "contradicting" | "missing"),
                             summary, excerpt
├── recommended_actions[]    each with action_type
                             ("immediate_corrective" | "long_term_corrective" |
                              "preventive" | "monitoring" | "procedure_update" |
                              "engineering_evaluation"),
                             priority ("critical" | "high" | "medium" | "low"),
                             target_component_id, rationale
├── barrier_analysis         summary of which barriers held, degraded, or failed
├── ccf_summary              common-cause failure assessment (injected post-synthesis)
├── human_performance_assessment optional; injected when Ishikawa ran and HP factors found
├── unresolved_gaps[]        specific evidence items that, if obtained, would close open questions
├── analyst_review
│   ├── decision_required (bool)
│   ├── questions_to_resolve[] open questions the analyst must address before writeback
│   └── writeback_recommendation "hold_until_review" | "ready_if_accepted"
└── validation_status
    ├── schema_valid (bool)
    ├── all_claims_cited (bool)
    ├── passed_minimum_evidence_gate (bool)
    └── fallback_used (bool)    true when LLM synthesis failed and rule-based fallback was used
```

If `fallback_used = true` the card was generated by the deterministic fallback template rather than the LLM synthesizer. The safety post-processing (barrier injection, CCF summary, human performance) always runs regardless of which path was taken.

2.7.2 run_manifest — audit and review record

The `run_manifest` is produced at Step 6i (§5.12). It is the artifact the reviewer uses to assess run quality, data completeness, and readiness for writeback. Key sections:

Section	Contents
<code>run_id</code> , <code>completed_at</code> , <code>input_refs</code>	Run identity and input data inventory
<code>pipeline_config</code>	Engine version, metamodel compliance level, top-k parameters, Ishikawa enabled flag, scoring evolution summary (pre- vs post-refine leader), scope revision state
<code>artifacts{}</code>	Per-artifact presence and health summary: counts (evidence items, pattern count, candidate count, barrier count, etc.) for every intermediate artifact
<code>coverage_summary</code>	Per-source data coverage status (<code>complete</code> / <code>partial</code> / <code>missing</code>); includes paired-data checks (SOE + protection logic)
<code>sensitivity_table</code>	For each missing or degraded source: whether supplying it could change the top-ranked candidate
<code>analyst_attention_flags[]</code>	All flags from §2.5.1, plus a sensitivity flag if any ranking change is possible
<code>review_hooks</code>	<code>next_step</code> , <code>writeback_ready</code> , <code>requires_human_review</code> , <code>analyst_decisions_required[]</code> (see §2.5.2)
<code>scope_revision_summary</code>	Active scope version, number of accepted revisions, latest analyst decision
<code>scope_expansion_summary</code>	Total expansion signals, pending analyst decisions
<code>pipeline_health</code>	Aggregate health status (<code>green</code> / <code>yellow</code> / <code>red</code>) and issues list
<code>stage_health</code>	Per-stage (A–G) health status
<code>similar_event_list</code>	Plant / fleet / industry match counts and degraded tier list (summary only; full list in intermediate artifact)
<code>signal_lessons_learned</code>	Matched pattern count, novel pattern flag, input sources used
<code>decision_trail</code>	Per-candidate scoring lineage for traceability
<code>analyst_checkpoints</code>	Ordered list of review checkpoints with status
<code>ap913_completeness</code>	AP-913 framework coverage assessment
<code>validation</code>	Input validation result, output validation result, optional artifact failures

2.7.3 Intermediate artifacts

These are persisted at each pipeline step and form the evidence trail behind the card. The analyst and data analyst use them to investigate specific sub-scores or trace a surprising ranking.

Artifact	Saved by step	Key contents
<code>kg_context.json</code>	Step 1	Equipment model neighborhood, past events, governance status, out-of-boundary anomalies
<code>signal_evidence.json</code>	Step 2a	Augmented anomaly list, propagation chains, chain warnings
<code>tskr_patterns.json</code>	Step 2b	Per-FM temporal patterns with Allen relations, onset lags, recurrence profiles, novel flags
<code>allen_relation_map</code>	Step 5b (§5.8)	Event-wide signal timeline; in-memory only — not persisted as a standalone file; embedded in manifest finalization
<code>causality_candidates_pre_refine.json</code>	Step 5 / §5.8	Pre-evidence-refinement candidate list; enables scoring evolution comparison

Artifact	Saved by step	Key contents
causality_candidates.json	Step 5 / §5.8	Final ranked candidates with composite scores, sub-scores, hard gate outcomes, causal categories
evidence_bundle.json	Step 5a / §5.8	Chroma-retrieved document chunks linked to candidates
ishikawa_matrix.json	Step 6a / §5.12 (optional)	Cause-and-effect structure; human performance and organizational factor branches
barrier_analysis.json	Step 6b / §5.12	Barrier states and degradation assessment
similar_event_list.json	Step 6c / §5.12	Full plant / fleet / industry match list with confidence weights
signal_lessons_learned.json	Step 6i / §5.12	Matched vs novel TSKR pattern summary
run_context.json	Steps 0, 4, 6h (§5.12)	Versioned scope, pipeline runtime flags, input availability map
workflow_dispatch.json	Step 6i / §5.12 (when enabled)	External review queue dispatch record with transport status
scoring_evolution.json	Step 6i / §5.12	Compact pre/post-refine rank comparison

`cmms_context.json` is persisted as an optional artifact when CMMS data was supplied and is flagged in `run_manifest.artifacts.cmms_context` .

2.7.4 What `decision_status` means

The `rca_card.executive_summary.decision_status` field is the most direct readout of the pipeline's confidence in the result:

Value	Meaning	Typical next step
"candidate_ready"	A primary hypothesis exists, all claims are cited, and evidence and severity gates passed	Read <code>analyst_attention_flags</code> ; if none are blocking, proceed to writeback review
"review_required"	A primary hypothesis exists but at least one condition is unresolved (uncited claim, open question, evidence gate not passed)	Resolve <code>analyst_review.questions_to_resolve</code> ; consider data backfill (§2.5.4)
"insufficient_evidence"	No candidate reached the minimum evidence threshold	Investigate data gaps in <code>coverage_summary</code> ; consider scope expansion (§2.5.3)

3. Data Structures: Knowledge Graph (Neo4j) and Chroma (Vector Store)

The pipeline relies on two fundamentally different persistence stores. The **Knowledge Graph** holds the structured equipment model — every component, failure mode, and safety function the pipeline is allowed to reason about. The **Chroma vector store** holds the documentary record — condition reports, work orders, procedures, and OE documents the pipeline retrieves as evidence. They answer different questions and are populated and maintained by different people.

Neither is modified by the pipeline itself during a run. Both are read-only inputs to any given RCA execution.

3.1 Knowledge Graph (Neo4j)

What it contains

The KG is the typed hypothesis space for the pipeline. It is a property graph stored in Neo4j and queried at run time by `Neo4jKGContextBuilder` to produce the `kg_context.json` snapshot that the orchestrator uses. The graph covers:

Node type	What it represents	Key properties
element_usage / element_definition	Physical components and their design definitions	component_id , component_label , component_type , maximo_floc , sap_equipment_id
failure_mode	Named failure modes applicable to a component	fm_id , name , superclass , expected_latency_min/max_hours , expected_symptoms , rpn
safety_function	Plant safety functions a component performs or supports	sf_id , sf_name , sf_category
monitored_variable	Process tags and sensor signals attached to a component	variable_id , monitored tag list
abnormal_event / oe_document	Past plant events and operating experience documents stored in the graph	event_id , severity , fm_id , resolved , priority_score
document	References to procedures, engineering documents, and bulletins	doc_id , doc_type , title , created_at

Edges encode relationships: `APPLIES_TO` (failure mode → component), `MONITORS / MEASURES` (signal → component), `has_part_usage` (containment hierarchy), `connects_port` (physical connectivity), `PERFORMS / SUPPORTS / ENABLES` (safety function links), `MAY_CAUSE / CONFIRMED_CAUSE / RELATED_TO` (past event causal links).

How data is structured at the pipeline boundary

The orchestrator does not query Neo4j directly during a run. `Neo4jKGContextBuilder.build()` issues a set of Cypher queries scoped to the event's `asset_id` and a configurable neighborhood (default: `max_hops = 2`) and serializes the result into a `kg_context.json` snapshot. This snapshot is the only KG input to the orchestrator. Its top-level fields:

Field	Contents
<code>components[]</code>	Components in the event neighborhood with their type, FLOC/SAP IDs, and monitored variable list
<code>failure_modes[]</code>	Failure modes applicable to those components with latency bounds, expected symptoms, and RPN
<code>safety_functions[]</code>	Safety functions associated with the neighborhood
<code>past_events[]</code>	Historical events from the graph, tagged with precursor window tier after <code>_enrich_past_events_temporal_metadata</code>
<code>upstream_paths[]</code>	Graph paths from the seed component to upstream components with path strength scores
<code>documents[]</code>	Document references from KG nodes (not the full text — that lives in Chroma)
<code>out_of_boundary_anomalies[]</code>	Signals detected outside the approved scope boundary (populated at Step 4)
<code>seed_context</code>	Asset IDs, monitored variables, and per-component past-event index used internally
<code>kg_snapshot_version</code>	Version string tied to KG software version and <code>last_modified</code> timestamp — used for replayability

The query scope is controlled by `KGContextBuilderConfig`: `max_hops` , `max_past_events` (default 10), `max_documents` (default 20), `doc_window_days_before/after` (default 90/7 days).

Who populates and maintains the KG

Role	Responsibility
Plant configuration engineers	Create and maintain component hierarchy, FLOC/SAP IDs, monitored variable assignments, and connectivity edges in Neo4j. This is the authoritative source of equipment topology.
Reliability / RCM engineers	Define and maintain failure modes: applicability to components, latency bounds, expected symptoms, RPN values. Gaps here directly limit the pipeline’s hypothesis space.
Plant event records staff	Enter past events (<code>abnormal_event</code> nodes) and their causal links (<code>MAY_CAUSE</code> , <code>CONFIRMED_CAUSE</code>) after each RCA or CR closure. These feed the recurrence and pattern-recognition analyses.

Role	Responsibility
Document control	Maintain document reference nodes pointing to procedures, engineering evaluations, and bulletins. Full text lives in Chroma; the KG holds metadata and link pointers only.
KG governance process	Periodically validates graph health (minimum failure mode count, stale snapshot detection). The pipeline reports governance status in <code>run_manifest.kg_governance</code> ; a “red” state can trigger a hard-abort (see §2.5.2).

The pipeline does not write to the KG during a run. Scope revision decisions (§2.5.3) update `run_context` only; they do not alter the graph.

3.2 Chroma (vector store)

What it contains

Chroma holds pre-indexed document chunks from the plant’s documentary record. Documents are split into chunks, embedded, and stored with structured metadata so the pipeline can retrieve the most relevant passages for each candidate during evidence retrieval (Step 5a / §5.8).

Document types indexed:

Type	Source	Priority weight
CR — Condition Report	Plant CMMS (e.g. Maximo)	1.00
WO — Work Order	Plant CMMS	0.95
ECA — Engineering Condition Assessment	Plant engineering	0.92
RCA — Root Cause Analysis report	Plant RCA program	0.90
ECR — Engineering Change Request	Plant engineering	0.85
FMEA — Failure mode and effects analysis	Plant engineering	0.80
SOP — Standard operating procedure	Plant operations	0.75
OE — Operating Experience	INPO, NRC, EPRI, owner-group	0.70
MANUAL — Vendor/maintenance manual	Vendor documentation	0.60
BULLETIN — Vendor or regulatory bulletin	Vendor / NRC	0.55

Each stored chunk carries metadata the retriever uses to filter and score:

Metadata field	Purpose
<code>record_id</code> , <code>doc_id</code> , <code>doc_type</code> , <code>chunk_index</code>	Identity and provenance
<code>authority_level</code>	<code>mandatory</code> / <code>guidance</code> / <code>informational</code> — affects evidence gate scoring
<code>primary_component_id</code> , <code>component_ids[]</code>	Enables component-scoped filtering during retrieval
<code>failure_mode_refs[]</code> , <code>causal_statements_text</code>	Pre-extracted causal content for scoring
<code>finding_status</code>	<code>confirmed</code> / <code>preliminary</code> / <code>fleet_experience</code> / <code>observational</code>
<code>dominant_temporal_relation</code>	Pre-extracted temporal signal (e.g. “precedes”, “during”)
<code>issuing_body</code> , <code>oe_number</code> , <code>plant_scope</code>	OE-specific provenance fields
<code>ca_as_found_condition</code> , <code>ca_as_left_condition</code>	Work order condition assessment (flattened from WO)

How retrieval works

`ChromaEvidenceRetriever` (Step 5.7) issues multiple typed queries against the store, one set per candidate plus event-level and component-level queries. Each query carries a `query_type` , a `weight` , and a `query_intent` (“support” or “contradiction”):

Query type	Weight	Intent
Candidate hypothesis support	1.00	Find passages that support this candidate's failure mode
Candidate contradiction check	0.70	Find passages that contradict this candidate
Failure mode context	variable	Broaden to the FM description
Component context	variable	Broaden to the component's documented history
Out-of-boundary anomaly	variable	Retrieve evidence for signals outside the scope boundary
Operational context / OE	variable	Fleet-wide and procedure context

Retrieval is **hybrid**: dense vector similarity (Ollama embeddings) combined with BM25 keyword matching. Each hit receives a composite score — `dense_score × doc_type_priority_weight × query_weight` — and results are sorted by that score. When the collection was loaded from disk rather than ingested in-process, BM25 is unavailable and retrieval degrades to dense-only; this is flagged in `evidence_bundle.provenance.quality_warning` .

Retrieved chunks are deduplicated, re-ranked by `evidence_score` , and assigned a `support_role` (`supporting` / `contradicting` / `contextual` / `missing`). The final `evidence_bundle.json` carries all results plus a `candidate_evidence_summary` that aggregates hit counts and best scores per candidate.

Who populates and maintains Chroma

Role	Responsibility
Document processing pipeline	Parses source documents (CR, WO, ECA, RCA, OE, manuals) into chunks, extracts metadata fields (causal statements, temporal references, finding status), and ingests them via <code>ChromaRecordStore.upsert_records()</code> . New documents are added incrementally; existing records are updated by re-upserting with the same <code>record_id</code> .
Plant CMMS administrators	Keep condition reports and work orders current. Closed CRs with confirmed root causes are the highest-value documents for the pipeline — they directly inform <code>causal_statements_text</code> and <code>finding_status = “confirmed”</code> .
Document control / OE program	Ingest NRC LERs, INPO SOERs, and EPRI reports as <code>OE</code> type documents. These feed the industry-tier evidence retrieval and the fleet OE similar event search (§2.6).
Pipeline engineers	Control <code>EvidenceRetrieverConfig</code> : <code>top_k_total</code> (default 10), <code>score_threshold</code> , <code>doc_type_priority</code> weights. Changes to these parameters directly affect what evidence reaches the RCA card.

The pipeline does not write to Chroma during a run. Pre-supplied `evidence_bundle.json` files (test/fixture mode) bypass retrieval entirely; the retriever’s live query path is skipped when `evidence_bundle` is passed directly to `run()` .

3.3 How KG and Chroma work together

The two stores serve complementary roles in producing the composite candidate score:

Question	Answered by	Pipeline artifact
Which failure modes are applicable to this component?	KG (<code>failure_modes[]</code>)	<code>causality_candidates</code> — structural sub-score
What is the expected degradation timescale for this FM?	KG (<code>expected_latency_min/max_hours</code>)	Temporal alignment in <code>tskr_patterns</code>
Has this component failed this way before at this plant?	KG (<code>past_events[]</code>) + CMMS	<code>similar_event_list</code> , recurrence sub-score
What did prior CRs and RCA reports say about this failure mode?	Chroma (CR, WO, RCA, ECA chunks)	<code>evidence_bundle</code> — evidence sub-score
Do vendor manuals or OE bulletins describe this failure mechanism?	Chroma (MANUAL, BULLETIN, OE chunks)	<code>evidence_bundle</code> — contextual citations in <code>rca_card</code>
Are there regulatory or procedure-level constraints relevant to this candidate?	Chroma (SOP, mandatory-authority chunks)	Evidence gate in <code>causality_candidates</code> post-refine

A failure mode that has no KG node is invisible to the pipeline regardless of how much evidence exists in Chroma. Conversely, a well-defined KG failure mode with no supporting Chroma documents will score low on the evidence dimension but will not be eliminated by a hard gate unless it also fails the timeline or barrier gates. **Maintaining both stores is necessary for the pipeline to produce well-supported results.**

5. Workflow Steps in Detail (System Engineer View)

This section provides step-by-step detail for each of the seven workflow stages introduced in §2.2. The audience is system engineers and data analysts who need to understand what the pipeline does at each stage, which data it depends on, and what artifacts it produces or transforms. Internal implementation details (function names, pseudocode, class hierarchies) are confined to code-anchor comments and appendices.

Note on execution order: §5 subsection numbering aligns with the §2.2 step numbering (Steps 0–6). A small number of operations execute at a different position than their conceptual step suggests — most notably, the Allen relation map (Step 2c conceptually) runs after Chroma evidence retrieval in the actual pipeline sequence. The cross-walk table below is the authoritative reference for these cases.

§2.2 step → §5 subsection cross-walk

§2.2 Step	§5 subsections	Notes
Step 0 — Initialize	§5.1	Run ID, input validation, run context, input guards, optional PM auto-build
Step 1 — Equipment model and data context	§5.2, §5.3	KG context build, governance, CMMS adapter, past-event enrichment (§5.2); signal evidence assembly (§5.3)
Step 2 — Temporal analysis and OE search (overview)	§5.4	Intro to all four sub-steps; execution order notes
Step 2a — Equipment history temporal search	§5.4 / Step 2a	Past-event pool grouped by component and FM; feeds recurrence pool for Step 2b
Step 2b — TSKR per-FM temporal scoring	§5.4 / Step 2b	Onset lag, anomaly, chain, and history scores blended into per-FM confidence ; tskr_patterns output
Step 2c — Allen relation map	§5.4 / Step 2c	Signal-interval classification; executes after Chroma retrieval in run order
Step 2d — Similar event identification	§5.4 / Step 2d	Plant-tier scoring + fleet/industry adapter query; executes before RCA card synthesis
Step 3 — Documentary pattern recognition	§5.5 / Step 3	Recurrence profiling (count, trend, unresolved) from kg_context.past_events ; embedded in TSKR run
Step 3.5 — Signal pattern recognition	§5.5 / Step 3.5	Matched vs novel classification (signal_lessons_learned); runs at manifest finalization
Step 4 — Candidate generation and initial ranking	§5.6	Five-dimension scoring, hard gates, chain position, scope boundary filter — all in one section
Step 5 — Evidence retrieval and refinement	§5.8	Chroma retrieval (5a); Allen map build (5b); evidence-informed re-scoring (5c); auto-reentry (5d) — all in one section
Step 6 — Conclusion and finalization	§5.12	Ishikawa (6a); barrier (6b); similar event list (6c); RCA card synthesis (6d); attention flags (6e); output validation (6f); Chroma archive (6g); scope expansion signals (6h); run manifest + return (6i) — all in one section

Appendix cross-reference: **A** — pattern-recognition mechanisms; **B** — LLM vs deterministic paths; **C** — similar-event query and tiers.

Each subsection uses the following format: **Purpose** — what the step does and why it matters; **Inputs** — artifacts and parameters consumed; **How it works** — plain-language description of the mechanism; **Outputs** — artifacts produced; **Parameters and defaults** — key configuration values; **What can go wrong** — failure modes and how they surface.

5.1 — Step 0: Initialize

Purpose Assign a unique identity to the run, verify that all supplied inputs meet the expected schema, and freeze the initial investigation scope — so that every downstream step operates on a consistent, validated dataset from a known starting point.

Inputs

Artifact	Required?	Description
event	Required	Equipment event being investigated (asset ID, timestamps, event type)
telemetry_summary	Required	Pre-processed anomaly records from the plant historian
operational_context	Optional	Shift narrative, recent alarms, and operational state at time of event
alarm_log	Optional	Alarm activations during and around the event window
soe_log	Optional	Sequence-of-events point records
protection_logic_context	Optional	Actuation logic and safety function state
configuration_change_records	Optional	Recent plant configuration changes
pm_compliance	Optional	Pre-built preventive maintenance compliance record; auto-built from CMMS export rows if not supplied
environmental_monitoring	Optional	Ambient conditions at event time; influences data coverage summary and Category F scoring
vendor_supply_chain_records	Optional	Lot/batch records; influences data coverage summary and Category C/K scoring
training_records	Optional	Personnel qualification records; influences data coverage summary and Category G scoring

How it works

The pipeline executes five actions in sequence:

- 1. Assign run identity.** A unique `run_id` is generated and a `run_status` record is immediately written to the artifact store with `run_complete: False`. This record is updated to `True` only at the end of Step 6 — if it remains `False`, the run did not complete.
- 2. Build PM compliance (if not supplied).** If no `pm_compliance` artifact was passed in, the pipeline attempts to construct one automatically from CMMS export rows in `operational_context`. Three modes are available: `auto` (default — build if absent), `off` (skip always), and `force` (re-run the build step even if no CMMS rows were found, but note: if `pm_compliance` was already supplied as an input, `force` has no effect — the guard is `if pm_compliance is None`). The look-back window defaults to **730 days**. PM compliance is treated as optional: if the build fails or the data are unavailable, the failure is recorded in `run_context.pipeline_runtime.pm_compliance` and the run continues.
- 3. Validate all inputs.** The four core artifacts (`event`, `telemetry_summary`, `operational_context`, `pm_compliance`) are checked against their JSON schemas. Required artifacts (`event`, `telemetry_summary`) must pass or the run stops. Optional artifacts that fail validation are logged in `optional_artifact_failures` and the run continues without them. A cross-check confirms that the `asset_id` in `event` matches the `asset_id` in `telemetry_summary`; mismatches are flagged as input guard warnings.
- 4. Build run context and initial scope.** The `run_context` record is assembled. It contains the run identity, a copy of the validated input metadata, and — critically — the initial investigation scope (scope version 0). The scope is seeded from the alarm log, SOE log, and operational context: systems in scope come from recent alarms; components in scope come from SOE entries and any CMMS records already available. This initial scope is **open** — it does not filter any candidates at this stage. Scope filtering only activates at Step 4 if an analyst has approved a scope revision from a prior run (see §5.6 and §2.5.3).
- 5. Enforce input guard policy.** After `run_context` is assembled, input guard flags are evaluated. Under strict configuration (`strict_input_guard_enforcement=True` or `input_guard_hard_stop_on_any_flag=True`), any flagged condition — including `asset_id` mismatch, missing required cross-checks, or output directory not writable — causes the orchestrator to raise immediately with `run_status.aborted=True`. This abort happens **before the KG build begins**. Under permissive configuration (the default), flags are recorded in `run_context` and the run continues.

Outputs

Artifact	Description
run_context	Run identity, scope version 0 (open), input guard flags, PM build metadata
input_validation	Per-artifact pass/fail results; <code>optional_artifact_failures</code> list

Parameters and defaults

Parameter	Default	Effect
pm_compliance_build_mode	auto	Controls whether PM compliance is auto-built (auto), skipped (off), or always rebuilt (force)
pm_compliance_look_back_window_days	730	How far back the PM auto-build searches for maintenance records
stop_on_validation_error	True	If True , a required-artifact validation failure aborts the run immediately

What can go wrong

- **asset_id mismatch** between event and telemetry_summary : flagged as an input guard warning; depending on policy, may stop the run or allow it to continue with a flag in run_context .
- **PM auto-build fails** (no CMMS export rows, adapter unavailable): failure is recorded in run_context.pipeline_runtime.pm_compliance with build_succeeded: False ; run continues but the PM compliance dimension in Step 4 scoring will be unavailable.
- **Input guard policy fires**: if strict enforcement is active and any input guard flag is set, the run aborts at Step 0 action 5 before the KG build begins. Check run_status.aborted and run_context.pipeline_runtime.input_guard_policy for the specific flag that triggered the stop.
- **Required artifact fails validation**: run stops immediately; the input_validation artifact records which field or schema rule was violated.
- **Scope seeds missing** (no alarm log, no SOE log, no operational context): scope version 0 is initialized with empty system and component lists; the pipeline still runs but Step 4 scope filtering (if active) will have nothing to filter against.

5.2 — Step 1 (part 1): Equipment model and data context

Purpose Materialize the typed hypothesis space — the set of components, failure modes, topology relationships, safety functions, and past events — that every downstream scoring and pattern-recognition step will reference. This step also enriches the hypothesis space with live CMMS records when available and normalizes all past-event timestamps relative to the triggering event.

Inputs

Artifact	Required?	Description
event	Required	Scopes the KG query: asset ID, timestamps, event type
telemetry_summary	Required	Used by the KG builder to link signal tags to graph entities
run_context	Required	Run identity and scope version produced by Step 0
operational_context	Optional	Shift and alarm data; used by the KG builder if available
pm_compliance	Optional	PM record from Step 0; used by the KG builder to assess maintenance posture
kg_context	Optional	Pre-built KG snapshot; if supplied, the live Neo4j query is skipped entirely (checkpoint-resume)

How it works

The pipeline executes four actions in sequence:

1. **Build the KG context.** If a pre-built kg_context is not supplied, the pipeline queries the Neo4j knowledge graph to retrieve: the asset's component tree (up to max_hops hops from the seed component), associated failure modes, topology and dependency edges, safety function assignments, and the document and past-event neighborhoods (CRs, WOs, ECAs, RCAs within the look-back window). The result is the typed hypothesis space for the run — all candidate failure modes and components that the causality engine will consider come from this artifact. If a pre-built snapshot is supplied, this query is skipped and the snapshot is used as-is.
2. **Assess KG governance.** The KG artifact is evaluated for data quality and coverage: are required node types populated, are relationships well-formed, are key attributes present? The result is a kg_governance record with a status of green , amber , or red . A **red** status means the KG is missing data critical to reliable scoring. By default (hard_abort_on_kg_red_state: True), a red status aborts the run immediately and records the specific issues in the manifest. An amber status allows the run to continue with a flag that appears in the final RCA card.
3. **Augment with CMMS records (if configured).** If a CMMS adapter is configured, the pipeline fetches the event-scoped CR and WO records from the CMMS system and merges them into kg_context in two ways: (a) as documents added to the document neighborhood (alongside KG-sourced

CRs and WOs), and (b) as past events added to the past-event pool (up to `cmms_past_event_injection_max` records). The merged `kg_context` is re-validated before proceeding. If the CMMS fetch fails for any reason, the failure is logged and the run continues with the KG-only context.

4. **Enrich past-event temporal metadata.** All past events in `kg_context.past_events` are enriched with normalized time relationships relative to the triggering event: days before the event, which precursor window tier they fall in, and whether they are within the active look-back window. This normalization is what Step 2a uses to filter past events by precursor window.

Outputs

Artifact	Description
kg_context	Components, failure modes, topology, safety functions, document neighborhood, past-event pool — enriched with CMMS records and temporal metadata
kg_governance	Data quality assessment (<code>green</code> / <code>amber</code> / <code>red</code>) with issue list; used by attention flags in Step 6
cmms_context	Raw CMMS CR/WO fetch result (only present if CMMS adapter is configured and fetch succeeds)

Parameters and defaults

Parameter	Default	Effect
max_hops	2	How many relationship hops from the seed component the KG query traverses
max_past_events	10	Maximum past events returned from the KG (before CMMS augmentation)
max_documents	20	Maximum documents (CRs, WOs, ECAs, etc.) in the document neighborhood
doc_window_days_before	90	Document look-back window before the event timestamp
doc_window_days_after	7	Document look-forward window after the event timestamp
cmms_past_event_injection_max	12	Maximum CMMS past events merged into <code>kg_context.past_events</code>
hard_abort_on_kg_red_state	True	If <code>True</code> , a red KG governance status aborts the run

What can go wrong

- **Red KG governance state:** run aborts (by default); the `kg_governance` record lists the specific missing nodes or broken relationships. Fix the KG data and re-run, or set `hard_abort_on_kg_red_state: False` to allow the run with reduced reliability.
- **CMMS fetch fails:** logged and skipped; run continues with KG-only past events and documents. The `run_manifest` will flag that CMMS augmentation was unavailable.
- **KG context is empty or minimal:** if the asset has few or no components, failure modes, or past events in the KG, downstream scoring will produce very low-confidence candidates. The `kg_governance` `amber/red` status is the primary signal for this condition.

5.3 — Step 1 (part 2): Signal evidence assembly

Purpose Build a unified, run-scoped record that maps each telemetry signal to its associated KG component, anomaly window, and propagation chain. This gives Steps 2–5 a structured view of which signals anomalized, in which order, and how they relate to components — rather than treating the raw telemetry summary as an unstructured list of readings.

Inputs

Artifact	Required?	Description
event	Required	Event identity and timestamps; anchors the signal time window
telemetry_summary	Required	Pre-processed anomaly records from the plant historian
kg_context	Required	Component and tag mapping from Step 1 (part 1); used to link signal tags to graph entities
signal_evidence	Optional	Pre-built signal evidence; if supplied, the build step is skipped (checkpoint-resume)

How it works

Each anomaly record in `telemetry_summary` carries a signal tag (e.g., a historian point ID). The builder resolves each tag to its corresponding component node in the KG and constructs a scored propagation chain — an ordered sequence of signal anomalizations that describes how the disturbance spread across components. When a plant historian adapter is configured, the builder can enrich the anomaly records with additional context (trend direction, rate of change, exceedance duration) directly from the historian. Without a historian adapter, the builder uses the pre-processed records in `telemetry_summary` as-is.

The output is the `signal_evidence` artifact: a structured record containing, for each signal, its component assignment, anomaly window (start, peak, end), propagation chain position, and any historian-sourced enrichment. If a signal tag cannot be resolved to a KG component, it is retained in `signal_evidence` with a null component link and flagged for the out-of-boundary attention check in Step 6.

Outputs

Artifact	Description
signal_evidence	Per-signal anomaly records with component links, propagation chains, and historian enrichment where available

Parameters and defaults

There are no analyst-configurable parameters for this step. The historian adapter (if any) is configured at the orchestrator level; its absence does not cause a validation failure as long as `telemetry_summary` provides valid anomaly records.

What can go wrong

- **Telemetry signals not resolvable to KG components:** signals are retained with null component links; downstream TSKR scoring will not be able to associate those signals with failure modes. The out-of-boundary attention flag in Step 6 will surface these gaps.
- **No historian adapter configured:** `signal_evidence` is built from `telemetry_summary` alone; historian enrichment fields will be absent. This is the normal operating mode in fixture/test environments.
- **telemetry_summary is empty or missing:** `signal_evidence` will be a minimal empty structure; TSKR scoring (Step 2b) will produce zero temporal support for all failure modes, and all TSKR-dependent causal categories will be flagged as data-limited in the manifest.

5.4 — Step 2: Temporal analysis and OE search

Step 2 establishes the full temporal and operating-experience picture before any hypothesis ranking. It has four sub-steps: (2a) filter the past-event pool by precursor window; (2b) score each failure mode on observed signal timing and recurrence history; (2c) classify every signal interval relative to the triggering event using Allen’s interval algebra; and (2d) identify similar past events at plant, fleet, and industry level. Sub-steps 2b and 2c share the same `tskr_patterns` and `allen_relation_map` artifacts that Step 4 and Step 5 consume directly.

Execution note: Sub-steps 2a and 2b execute in pipeline order here (after Step 1, before Step 4). Sub-step 2c (Allen map) executes after Chroma retrieval in the actual run sequence because it is used by evidence refinement, not by candidate generation. Sub-step 2d (similar events) executes just before RCA card synthesis. Both are described here for conceptual coherence; their execution positions are noted in each sub-section.

Step 2a — Equipment history temporal search

Purpose Filter the past-event pool that was loaded into `kg_context` at Step 1 down to records relevant to the current investigation, grouped by component and failure mode. This filtered pool is the input to the recurrence analysis in Step 2b.

Inputs `kg_context.past_events` (enriched with temporal metadata at Step 1), `event` (timestamp as reference point).

How it works Each past event in `kg_context.past_events` carries normalized temporal metadata added at Step 1: how many days before the triggering event it occurred and whether it falls within the active look-back window. Step 2a groups these records by `component_id` and `fm_id` so that Step 2b can immediately retrieve the relevant subset for any given failure mode without scanning the full pool.

```
for each (component_id, fm_id) pair in kg_context.failure_modes:

    recurrence_pool = [
        e for e in kg_context.past_events
        if e.component_id == component_id
        OR e.fm_id == fm_id
    ]

    # each event already carries:
    #   days_before_event    - days before event.timestamp_start
    #   within_window        - True if days_before_event ≤ look_back_days
    #   resolved             - True / False from CR/WO finding_status
    #   root_cause_label     - from prior RCA or CR if available

recurrence_pool_index[component_id][fm_id] = recurrence_pool
```

Outputs In-memory recurrence pool index, used immediately by Step 2b. No separate artifact is persisted; the pool is embedded in the `tskr_patterns` scoring context.

What can go wrong If `kg_context.past_events` is empty (no historical records in KG and no CMMS adapter configured), every failure mode will have an empty recurrence pool. All `recurrence_count` values in Step 2b will be 0 and all `novel_pattern` flags will fire regardless of signal evidence.

Step 2b — TSKR per-FM temporal scoring

Purpose For every failure mode in the equipment model, compute a `confidence` score that captures both how well the observed signal timing supports that failure mode and how often it has recurred in the plant’s documentary history. The result — `tskr_patterns` — is the primary temporal input to Step 4 candidate ranking.

Inputs

Artifact	Required?
event	Required
telemetry_summary	Required
kg_context	Required
signal_evidence	Optional (improves tag-to-FM mapping)
alarm_log	Optional (onset timing)
soe_log	Optional (sequence ordering)
tskr_patterns	Optional (pre-built; skips this step if supplied)

How it works For each failure mode, four sub-scores are computed and blended:

- Anomaly score (weight 0.45):* counts and peak magnitudes of anomaly hits on signal tags associated with the failure mode. More hits and higher exceedances raise this score.
- Onset lag score (weight 0.30):* measures whether signals began anomalizing before the triggering event and whether the lead time matches the expected degradation profile. Signals within `simultaneous_epsilon_hours` (0.5 h) of the event are treated as simultaneous rather than strictly preceding.
- Propagation chain score (weight 0.10):* checks whether the sequence of signal anomalizations across components matches the failure mode’s expected propagation path.
- History score (weight 0.10):* derived from the recurrence pool built in Step 2a (see Step 3 details in §5.5). Base score by count: 0 events → 0.0; 1 → 0.35; 2–3 → 0.55; 4–6 → 0.70; >6 → 0.80. Bonuses: +0.15 if trend is increasing; +0.10 if any events are unresolved; +0.05 if most recent event was within 90 days.

```
for each failure_mode in kg_context.failure_modes:

    anomaly_score      = score_anomaly_hits(telemetry_summary, alarm_log, failure_mode)
    onset_score        = score_onset_lag(signal_evidence, failure_mode, event,
                                       epsilon_hours=0.5)

    chain_score        = score_propagation_chain(signal_evidence, failure_mode)
    history_score       = score_history(recurrence_pool_index[failure_mode]) # from 2a
    anomaly_count_score = score_anomaly_count(telemetry_summary, failure_mode)
    lag_consistency_score = score_lag_consistency(signal_evidence, failure_mode)

    # Six-term normalized weighted sum
    raw = (0.45 × anomaly_score      +
           0.30 × onset_score        +
           0.10 × chain_score        +
           0.10 × history_score      +
           0.15 × anomaly_count_score +
           0.10 × lag_consistency_score)
    confidence = raw / (0.45+0.30+0.10+0.10+0.15+0.10) # normalize to [0,1]

    # novel_pattern requires all three conditions:
    # (1) no recurrence in documentary record
    # (2) history score below the novelty floor
    # (3) no signal IDs resolved to this failure mode
    novel_pattern = (recurrence_count == 0
                     AND history_score < 0.20
                     AND NOT bool(matched_signal_ids))

    tskr_patterns[failure_mode] = {confidence, onset_lag_hours,
                                   recurrence_count, recurrence_trend,
                                   unresolved_recurrence_count, novel_pattern,
                                   history_score}
```

Note on novel_pattern : a failure mode that has matching signal IDs (even if first-occurrence in the documentary record) is **not** flagged as novel — the third condition ensures that a signal match alone is enough to suppress the flag. True novelty requires no documentary trace **and** no signal match.

If no TSKR scorer is configured, a synthetic empty `tskr_patterns` is returned with `summary.mode: "absent"` and confidence = 0 for all failure modes.

Outputs

Artifact	Key fields
tskr_patterns	Per-FM: confidence , onset_lag_hours , recurrence_count , recurrence_trend , unresolved_recurrence_count , novel_pattern , history_score
tskr_patterns.summary	n_supported_patterns (confidence ≥ 0.35), n_novel_patterns , mode ("deterministic_v1" when scored; "absent" when no scorer configured)

Parameters and defaults

Parameter	Default	Effect
simultaneous_epsilon_hours	0.5 h	Tolerance for "simultaneous" signal onset
min_confidence_for_support	0.35	Threshold for counting a pattern as supported
anomaly_weight	0.45	Weight in confidence blend
latency_weight	0.30	Weight in confidence blend
chain_weight	0.10	Weight in confidence blend
history_weight	0.10	Weight in confidence blend
anomaly_count_weight	0.15	Weight for anomaly count score in confidence blend (not displayed in sub-score breakdown)

Parameter	Default	Effect
lag_consistency_weight	0.10	Weight for lag consistency score in confidence blend (not displayed in sub-score breakdown)

What can go wrong

- **No TSKR scorer configured:** all confidence values are 0; temporal dimension is unavailable in Step 4 scoring.
- **Empty recurrence pool:** all `recurrence_count = 0` , all `novel_pattern = True` . Does not stop the run; reduces confidence in the history component.
- **Signal tags not resolved to failure modes:** anomaly and onset scores default to fallback (0.25 each). Verify tag-to-FM mappings in `kg_context` .

Step 2c — Allen relation map

Execution note: The Allen map is built in pipeline execution order *after* Chroma evidence retrieval (Step 5) — not immediately after TSKR. It is described here because it is conceptually part of the Step 2 temporal picture. The artifact is used by evidence refinement (`refine_with_evidence`) and scope expansion signal detection, not by candidate generation.

Purpose Classify every signal interval in the event window — telemetry anomalies, alarm activations, SOE transitions — relative to the triggering event using Allen’s interval algebra. The result establishes which signals preceded the event (potential causes) versus which followed it (potential consequences), providing a shared temporal ordering map for evidence refinement and scope expansion.

Inputs `event` , `telemetry_summary` , `alarm_log` (optional), `soe_log` (optional).

How it works Each signal is represented as an interval `[start, end]` . The triggering event’s `[timestamp_start, timestamp_end]` is the reference interval. Every signal interval is classified into one of five Allen relations by testing membership conditions in order (FOLLOWS → PRECEDES → CONTAINS → OVERLAPS → DURING):

Relation	Condition	Base score
PRECEDES	Signal ends before event starts (within epsilon)	0.75
FOLLOWS	Signal starts after event ends (within epsilon)	0.10
CONTAINS	Event interval is fully inside signal interval	0.85
DURING	Signal interval is fully inside event interval	0.30
OVERLAPS	Intervals partially overlap	0.90

The Allen base score for each matched signal is then blended into the matching candidate’s temporal sub-score using a raise-only rule: the existing temporal score can only increase, never decrease.

```
for each signal_interval in (telemetry anomalies + alarm activations + SOE events):

    relation = classify_allen(signal_interval, event_interval, epsilon_hours=0.5)
    base_score = RELATION_SCORES[relation] # from table above

    for each candidate that maps to this signal:
        old_temporal = candidate.temporal_score
        blended      = 0.75 × old_temporal + 0.25 × base_score
        new_temporal = max(old_temporal, blended) # raises only
        candidate.temporal_score = new_temporal
```

Outputs In-memory `allen_relation_map` (passed to `refine_with_evidence` and manifest finalization; not a standalone top-level return artifact).

Parameters and defaults

Parameter	Default	Effect
epsilon_hours	0.5 h	Tolerance for boundary conditions (e.g. signal ending exactly at event start)
ALLEN_ALPHA	0.25	Weight of the Allen base score in the temporal blend (1 – ALLEN_ALPHA = 0.75 for existing score)

What can go wrong

- **No alarm log or SOE log supplied:** only telemetry anomaly intervals are classified; alarm and SOE ordering will be absent from the map. Scope expansion signals that rely on alarm ordering may not fire.
- **Timestamp inconsistencies** (signals with `start > end` , or NaT values): those intervals are skipped; a warning is logged. Ensure historian export and SOE timestamps are consistent before running.

Step 2d — Similar event identification

Execution note: The similar event list is assembled in pipeline execution order just before RCA card synthesis (after evidence refinement). It is described here because it is conceptually part of the Step 2 operating-experience picture. See §2.6 for full detail on the query logic and adapter contract.

Purpose Identify past events at three levels — plant, fleet, and industry — that are similar to the current event based on component, failure mode, event type, and actuation type. The resulting `similar_event_list` is provided to the RCA card synthesizer and appears in the final report as the "have we seen this before?" section.

Inputs `event` , `kg_context` (plant past events), `causality_candidates` (top candidates drive the query), `SimilarEventAdapter` (optional; enables fleet and industry tiers).

How it works

Plant tier: Each past event in `kg_context.past_events` is scored against the current event on five dimensions. The top `step2d_plant_top_n` events by score are returned.

```
for each past_event in kg_context.past_events:

    plant_score = (0.40 × component_match      +
                  0.25 × failure_mode_match    +
                  0.15 × event_type_match      +
                  0.10 × actuation_type_match  +
                  0.10 × precursor_window_overlap)

    similar_event_list (plant) = top step2d_plant_top_n events by plant_score
```

Fleet and industry tiers (if `SimilarEventAdapter` is configured):

```
fleet_matches    = adapter.query(level="fleet",
                                component_ids    = top candidate component IDs,
                                failure_mode_ids = top candidate FM IDs,
                                event_type       = event.event_type,
                                actuation_type   = event.actuation_type,
                                max_results      = step2d_query_top_n_candidates)

industry_matches = adapter.query(level="industry", ... )

similar_event_list = plant_matches + fleet_matches + industry_matches
```

Outputs `similar_event_list` — list of similar event records, each carrying `event_id` , `source_level` (plant / fleet / industry), `confidence_weight` , `summary` , `root_cause_label` , `resolution` , `lessons_learned_ref` , `contributing_categories` .

Parameters and defaults

Parameter	Default	Effect
step2d_plant_top_n	5	Maximum plant-tier matches returned
step2d_query_top_n_candidates	3	Number of top candidates whose component/FM IDs are used to query fleet and industry adapters

Parameter	Default	Effect
Plant-tier dimension weights	0.40 / 0.25 / 0.15 / 0.10 / 0.10	Component / FM / event_type / actuation_type / precursor_window

What can go wrong

- **No fleet/industry adapter configured:** only plant-tier matches are returned. The `similar_event_list` will be populated from `kg_context.past_events` only.
- **Plant past-event pool is sparse:** few or no plant matches; `similar_event_list` may be empty. The RCA card will note the absence but will not fail.
- **Adapter timeout or failure:** the failed tier is skipped and logged; plant-tier results are still returned. See §2.6 for the adapter error contract.

5.5 — Step 3 / 3.5: Pattern recognition

Step 3 and Step 3.5 both produce pattern classification outputs that characterize what the pipeline "already knows" about the current failure mode from prior plant experience and from the observed signal record. Neither step is a separate code block: the Step 3 documentary analysis is embedded in the TSKR scoring run (§5.4 / Step 2b), and Step 3.5 signal classification runs at manifest finalization (§5.18 / Step 6). They are described here as a pair because they answer a single SE question: *is this a known pattern or a novel one?*

Step 3 — Documentary pattern recognition

Execution note: This analysis runs inside TSKR scoring (Step 2b). Its outputs are already in `tskr_patterns` by the time Step 4 runs. This section describes the outputs and their interpretation.

Purpose For each failure mode, determine how well the plant's documentary record (past CRs, WOs, confirmed RCA findings) supports it as a known recurring issue. The recurrence profile answers: how many times has this failed before, is the rate increasing or decreasing, and are prior occurrences resolved?

Inputs `kg_context.past_events` (structured past events carrying `root_cause_label`, `finding_status`, `causal_explanation` from injected CRs/WOs/RCA reports), grouped into the recurrence pool by Step 2a.

How it works

for each (component_id, fm_id) pair:

```
pool = recurrence_pool_index[component_id][fm_id] # from Step 2a

recurrence_count = len(pool)
intervals_days = inter-event intervals sorted by date
unresolved_count = count where e.resolved == False
most_recent_days = days since most recent event in pool

# trend: ratio = second_half_mean / first_half_mean
# ratio < 1 means recent intervals are shorter → events accelerating → "increasing"
# ratio > 1 means recent intervals are longer → events slowing → "decreasing"
if len(intervals_days) < 3:
    trend = "insufficient_data"
else:
    ratio = mean(second_half) / mean(first_half)
    trend = "increasing" if ratio < 0.75 # events accelerating
           = "decreasing" if ratio > 1.33 # events slowing
           = "stable" otherwise

# history score
base = {0: 0.00, 1: 0.35, 2-3: 0.55, 4-6: 0.70, >6: 0.80}[recurrence_count]
history_score = base
               + 0.15 if trend == "increasing"
               + 0.10 if unresolved_count > 0
               + 0.05 if most_recent_days < 90
```

Outputs Recurrence fields on each `tskr_patterns` entry: `recurrence_count` , `recurrence_trend` , `unresolved_recurrence_count` , `most_recent_days_ago` , `history_score` , `novel_pattern` .

What to look for as an analyst

- A high `recurrence_count` with `trend = "increasing"` and `unresolved_count > 0` is a strong signal that the root cause has not been corrected. The attention flag `recurrence_match_quality` will fire if this combination is present in the top candidates.
- `novel_pattern = True` across all top candidates means the pipeline has no prior documentary trace for the failure modes it is most confident about. This warrants closer scrutiny of the evidence and a review of whether the plant's CR/WO history is fully loaded into the KG.

Step 3.5 — Signal pattern recognition

Execution note: This classification runs at manifest finalization (Step 6i / §5.12), after all scoring and synthesis are complete. It is described here for conceptual continuity with Step 3.

Purpose Classify each TSKR pattern as *matched* (the signal shape has a prior documentary trace) or *novel* (the signal pattern has never been recorded for this failure mode at this plant). This classification is assembled into `signal_lessons_learned` and stored in the run manifest as a structured lessons-learned record.

Inputs Finalized `tskr_patterns` (all per-FM fields including `recurrence_count` , `history_score` , `novel_pattern`).

How it works

```
for each pattern in tskr_patterns:

    if pattern.recurrence_count > 0 OR pattern.history_score >= threshold:
        classification = "matched"      # historical support found
    else:
        classification = "novel"        # no prior documentary trace

    signal_lessons_learned[pattern.fm_id] = {
        classification:      classification,
        novel_pattern:       pattern.novel_pattern,
        recurrence_count:    pattern.recurrence_count,
        recurrence_trend:    pattern.recurrence_trend,
        confidence:          pattern.confidence
    }
```

Outputs `signal_lessons_learned` — per-FM matched/novel classification; embedded in `run_manifest` . Not a standalone return artifact.

What to look for as an analyst

- Multiple "novel" entries across the top-ranked candidates suggest a genuinely new failure mode or a data gap (KG not populated for this asset). Cross-check against §3 (KG and Chroma coverage) before concluding novelty.
- A "matched" entry with high `recurrence_count` and unresolved prior events is the most actionable pattern: the pipeline has seen this before and it was not fixed.

5.6 — Step 4: Candidate generation and initial ranking

Purpose Generate all plausible failure-mode hypotheses from the equipment model, score each one on five evidence dimensions, eliminate physically or logically impossible candidates through hard gates, and produce a ranked shortlist for evidence retrieval. This is the first pass: scores are preliminary because Chroma evidence has not yet been retrieved. A scope boundary filter then removes any candidates outside the analyst-approved investigation boundary before evidence retrieval begins.

Inputs

Artifact	Required?	Description
event	Required	Event identity and reference interval
telemetry_summary	Required	Anomaly records for telemetry sub-score
kg_context	Required	Full equipment model: components, failure modes, topology, safety functions
tskr_patterns	Required	Per-FM temporal support and recurrence profiles from Steps 2b / 3
operational_context	Optional	Shift context; contributes to operating-point scoring
pm_compliance	Optional	PM posture; contributes to governance sub-score
run_context	Required	Scope version; drives the scope boundary filter
causality_candidates	Optional	Pre-built; skips candidate generation entirely (checkpoint-resume)

How it works

Candidate generation. Every failure mode in `kg_context` becomes a candidate. The engine does not generate hypotheses from free text — the hypothesis space is strictly bounded by the KG equipment model.

Five-dimension scoring. Each candidate receives a sub-score on five dimensions, which are blended into a single `composite_score` :

Dimension	Default weight	What it measures
Structural	0.30	KG topology fit: is this FM reachable from the affected component? CCF and operating-point modifiers applied here

Dimension	Default weight	What it measures
Temporal	0.20	TSKR confidence from <code>tskr_patterns</code> blended with Allen-based timing from <code>tskr_index</code> ; recurrence contribution also enters here
Telemetry	0.20	How well the observed anomaly pattern matches the failure mode's expected signal signature
Evidence	0.20	Initial estimate of documentary support; low at this stage (Chroma not yet queried)
Governance	0.10	PM compliance posture; weight can vary per failure mode superclass

Two additional adjustments are applied within the structural sub-score: a CCF delta (up to +0.10 for Category C common-cause contributors) and an operating-point delta (up to +0.12 for Category E environmental / operating-point contributors).

Hard elimination gates. After scoring, three binary gates are applied in sequence. Any candidate that fails a gate is moved to a documented ruled-out list with the gate name as the reason — it is retained in the output for audit purposes but is not carried forward to evidence retrieval. A dual-threshold check then filters candidates that pass all gates.

```
for each candidate:

    # Gate 1 – Physical plausibility
    if failure_mode is physically impossible given component state:
        ruled_out(reason="physically_impossible")
        continue

    # Gate 2 – Timeline consistency
    if candidate's causal timeline contradicts the observed signal sequence:
        ruled_out(reason="timeline_inconsistent")
        continue

    # Gate 3 – Barrier logic
    if a safety barrier confirmed as intact makes the failure mode impossible:
        ruled_out(reason="barrier_held")
        continue

    # Dual threshold check – BOTH conditions must be satisfied:
    #   (a) composite score meets the composite floor
    #   (b) evidence sub-score meets the evidence floor (meets_evidence_threshold flag)
    if composite_score < minimum_composite_threshold (0.30):
        filtered_out(reason="below_threshold")
        continue
    if NOT meets_evidence_threshold: # evidence sub-score < minimum_evidence_threshold
        filtered_out(reason="below_evidence_threshold")
        continue

    active_candidates.append(candidate)
```

Note: `minimum_evidence_threshold` (default 0.35) is the evidence sub-score floor for `meets_evidence_threshold` .

`minimum_pre_evidence_threshold` (default 0.10) is a separate, lower pre-evidence floor — a candidate that scores below 0.10 on the evidence dimension at Step 4 is immediately ineligible; `minimum_evidence_threshold` is checked after the evidence refinement pass in Step 5.

Chain position assignment. Each passing candidate is assigned a `chain_position` indicating its role in the causal sequence: `initiating` (temporal precedence indicates it triggered the event), `consequence` (temporal data shows it followed the trigger), or `contributing` (concurrent contributor without clear precedence).

Retention. The active candidate list is sorted by `composite_score` descending. Only the top `top_k_candidates` (default 5) are retained for evidence retrieval; any beyond that are moved to filtered-out with reason `"below_top_k"` .

```
composite_score = (0.30 × structural  +
                  0.20 × temporal    +
                  0.20 × telemetry    +
                  0.20 × evidence     +
                  0.10 × governance)
```

```
sort candidates by composite_score descending
active_candidates = top top_k_candidates
```

Event analogs. Alongside the failure-mode (FM) candidate pool, `generate()` also produces an `event_analogs` pool — past plant events whose profile closely matches the current event. Event analogs are scored and threshold-filtered independently; they do not compete for FM candidate slots and are not subject to the `top_k_candidates` cap. Both pools are returned in `causality_candidates` and consumed by the synthesizer and manifest.

Scope boundary filter. After `generate()` returns, the orchestrator applies the scope boundary filter. This is an **orchestrator-level operation** — the causality engine has no knowledge of scope; it operates only on the KG equipment model. On the first run, the scope is open (version 0) and no candidates are filtered. If the analyst has approved a scope revision (version ≥ 1) from a prior run, the orchestrator moves any candidate whose `component_id` is outside the approved component set to the ruled-out list with reason “`scope_filtered`” before evidence retrieval begins.

```
# Applied by orchestrator after causality_engine.generate() returns
if run_context.active_scope_version >= 1:
    approved_components = run_context.scope_management.approved_component_set

    for each candidate in active_candidates:
        if candidate.component_id NOT in approved_components:
            ruled_out(reason="scope_filtered",
                      scope_version=active_scope_version)

    run_context.pipeline_runtime.scope_filter.applied = True
else:
    run_context.pipeline_runtime.scope_filter.applied = False
```

Outputs

Artifact	Key fields
causality_candidates	candidates list (active FM hypotheses, scored, chain-positioned); event_analogs list (past-event profile matches); ruled_out list with reason codes; scoring_config echo; scope_filter metadata

Each FM candidate carries: `candidate_id`, `component_id`, `fm_id`, `composite_score`, `scores` (per-dimension), `hard_gates` (pass/fail per gate), `chain_position`, `chain_position_confidence`, `meets_evidence_threshold`.

Parameters and defaults

Parameter	Default	Effect
structural weight	0.30	Contribution of KG topology to composite score
temporal weight	0.20	Contribution of TSKR / Allen temporal support
telemetry weight	0.20	Contribution of anomaly signal fit
evidence weight	0.20	Contribution of documentary evidence (low at this stage)
governance weight	0.10	Contribution of PM compliance posture
minimum_composite_threshold	0.30	Candidates below this composite score are filtered out
minimum_pre_evidence_threshold	0.10	Minimum evidence sub-score at Step 4 entry; candidates below this floor are immediately ineligible
minimum_evidence_threshold	0.35	Minimum evidence sub-score after Step 5 refinement; candidates below this threshold have <code>meets_evidence_threshold=False</code> and fail the dual-threshold check

Parameter	Default	Effect
top_k_candidates	5	Maximum FM candidates carried forward to evidence retrieval (OrchestratorConfig default; causality engine internal default is 10 but is overridden at runtime)

What can go wrong

- **All candidates eliminated by hard gates:** typically indicates a data issue — either the KG’s safety function state is inconsistent, or the signal timestamps are unreliable. Check kg_governance (Step 1) and signal timestamps (Step 0 asset_id cross-check).
- **All candidates below composite threshold:** the equipment model may be incomplete (too few failure modes in KG) or all TSKR scores are zero (no scorer configured). Review kg_context coverage and TSKR scorer configuration.
- **Scope boundary filters too aggressively:** if a scope revision was approved with a narrow component set, legitimate candidates may be excluded. Re-run without the scope filter (scope version 0) to confirm, then revise the approved component set (see §2.5.3).
- **Chain position all “contributing”:** no temporal data strong enough to distinguish initiating vs consequence candidates. Supply soe_log and alarm_log to improve Allen and TSKR onset scoring.

5.8 — Step 5: Evidence retrieval and refinement

Step 5 takes the ranked candidate list from Step 4 and gives it a second, evidence-informed scoring pass. It runs in four sub-steps: (5a) retrieve relevant document passages from Chroma for each candidate; (5b) build the Allen relation map (this is where Step 2c conceptually executes); (5c) re-score every candidate using the retrieved evidence, Allen signals, and data coverage state; and (5d) optionally re-run the full Steps 1–5 sequence with a broader context if the first pass produced low-confidence results. The causality_candidates artifact produced by this step (v2, post-refine) is what the RCA card synthesis in Step 6 uses.

Step 5a — Chroma evidence retrieval

Purpose Query the Chroma vector store for document passages relevant to each active candidate, populating the evidence score dimension with real documentary support. Passages retrieved here are also cited directly in the final RCA card.

Inputs

Artifact	Required?
event	Required
kg_context	Required
causality_candidates (v1, from Step 4)	Required
operational_context	Optional
evidence_bundle	Optional (pre-built; skips retrieval if supplied)

How it works For each active candidate, the retriever issues two queries against the Chroma store: a *support* query (looking for passages consistent with the failure mode) and a *contradiction* query (looking for passages that argue against it). Each retrieved passage is scored multiplicatively:

$$\text{passage_score} = \text{dense_score} \times \text{doc_type_priority_weight} \times \text{query_weight}$$

Query weights: support = 1.00; contradiction = 0.70. Document type priority weights:

Doc type	Weight	Doc type	Weight
CR	1.00	ECR	0.85
WO	0.95	FMEA	0.80
ECA	0.92	SOP	0.75
RCA	0.90	OE	0.70

Doc type	Weight	Doc type	Weight
—	—	MANUAL	0.60
—	—	BULLETIN	0.55

The top `top_k_total` passages across all candidates and queries are retained. BM25 keyword matching is applied alongside dense vector similarity when available.

```
for each candidate in active_candidates:

    support_passages = chroma.query(
        text = failure_mode description + component context,
        filter = {component_id, fm_id},
        type = "support", weight = 1.00
    )
    contradiction_passages = chroma.query(
        text = negation of failure_mode description,
        filter = {component_id, fm_id},
        type = "contradiction", weight = 0.70
    )
    for each passage:
        passage_score = dense_score × doc_type_priority_weight × query_weight

evidence_bundle = top top_k_total passages ranked by passage_score
```

Outputs `evidence_bundle` — per-candidate passage records with scores, doc type, source reference, and query type (support / contradiction / contextual).

Parameters and defaults

Parameter	Default	Effect
<code>top_k_total</code>	10	Maximum passages retained across all candidates
<code>query_weight (support)</code>	1.00	Full weight for passages supporting a hypothesis
<code>query_weight (contradiction)</code>	0.70	Reduced weight for passages contradicting a hypothesis

What can go wrong

- **Chroma store empty or unindexed:** `evidence_bundle` will have zero passages; the evidence sub-score for all candidates stays at the Step 4 initial estimate. The run continues but the RCA card will cite no documents.
- **Pre-built `evidence_bundle` supplied:** retrieval is skipped entirely (checkpoint-resume). Useful for iterative runs where the document set has not changed.

Step 5b — Allen relation map (execution)

Conceptual note: This is where the Allen relation map described in Step 2c (§5.4) is actually built in the pipeline execution sequence — after Chroma retrieval, immediately before evidence-informed re-scoring. The map is passed directly to Step 5c and is also used by scope expansion signal detection in Step 6.

Purpose Classify every signal interval (telemetry anomalies, alarm activations, SOE transitions) relative to the triggering event using Allen’s interval algebra. The result tells Step 5c which signals preceded the event (potential causes) and which followed it (potential consequences), allowing the re-scoring pass to sharpen temporal sub-scores.

Inputs `event` , `telemetry_summary` , `alarm_log` (optional), `soe_log` (optional).

How it works Each signal interval `[start, end]` is compared to the event interval `[timestamp_start, timestamp_end]` using the five Allen relations (see Step 2c in §5.4 for the full relation table and base scores). The result is an in-memory `allen_relation_map` passed immediately to Step 5c.


```
allen_relation_map = {}

for each signal_interval in (telemetry + alarms + SOE):
    relation = classify_allen(signal_interval, event_interval, epsilon_hours=0.5)
    base_score = {PRECEDES: 0.75, FOLLOWS: 0.10, CONTAINS: 0.85,
                  DURING: 0.30, OVERLAPS: 0.90}[relation]
    allen_relation_map[signal_id] = {relation, base_score}
```

Outputs In-memory allen_relation_map (not a standalone persisted artifact; passed to Step 5c and Step 6 manifest finalization).

Step 5c — Evidence-informed re-scoring

Purpose Re-score every active candidate using the retrieved evidence passages, Allen relation signals, protection logic, and data coverage state. This is the second and final scoring pass — the composite_score values produced here are the ones that appear in the RCA card.

Inputs

Artifact	Required?
causality_candidates (v1)	Required
evidence_bundle	Required
signal_evidence	Required
allen_relation_map	Required (from Step 5b)
protection_logic_context	Optional
environmental_monitoring	Optional
vendor_supply_chain_records	Optional
training_records	Optional

How it works

Before re-scoring, the pre-refine state is snapshotted as causality_candidates_pre_refine . This snapshot is used by the rank-inversion attention flag (Step 6) and by auto-reentry (Step 5d) to detect whether re-scoring changed the ranking.

A data coverage summary is assembled from all available artifacts. This summary tells the engine which data dimensions were populated and which were absent, allowing it to weight sub-scores appropriately when data is sparse.

The re-scoring pass then updates each candidate:

```

# Snapshot before re-scoring
causality_candidates_pre_refine = copy of causality_candidates (v1)

# Coverage summary
coverage_summary = {
    kg_coverage:          kg_context completeness metrics,
    tskr_coverage:        tskr_patterns.summary,
    evidence_coverage:     evidence_bundle retrieval metrics,
    environmental:         "available" or "absent",
    vendor_supply_chain:   "available" or "absent",
    training_records:      "available" or "absent"
}

# Pass 1 – Evidence re-scoring loop (one candidate at a time)
for each candidate in causality_candidates (v1):

    support_score      = sum of passage scores where support_role == "supporting"
    contradiction_score = sum of passage scores where support_role == "contradicting"
    contextual_score    = sum of passage scores where support_role == "contextual"
    authority_weight    = doc_authority_multiplier(best_passage.authority_level)

    # Four-term evidence blend (clipped to [0, 1]):
    new_evidence = max(0, min(1,
        0.30 × prior_evidence_score
        + 0.55 × support_score × authority_weight
        + 0.15 × contextual_score
        - 0.45 × contradiction_score
    ))

    # If chain position evidence is available, blend doc-based and chain-based scores:
    # new_evidence = 0.70 × new_evidence_doc + 0.30 × new_evidence_chain

    # Refine temporal sub-score using Allen signals (raise-only)
    allen_score = allen_relation_map.get(candidate.component_signals)
    new_temporal = max(old_temporal, 0.75 × old_temporal + 0.25 × allen_score)

    # Category minimum evidence gate (4th gate – applied inside evidence loop)
    if new_evidence < category_minimum_evidence_threshold[candidate.causal_category]:
        candidate.category_evidence_blocked = True

    new_composite = (0.30 × structural +
        0.20 × new_temporal +
        0.20 × telemetry +
        0.20 × new_evidence +
        0.10 × governance)

# Pass 2 – Hard gate re-check loop (runs after all evidence scoring is complete)
for each candidate in causality_candidates (v1):
    apply_physical_plausibility_gate(candidate, kg_context)
    apply_timeline_consistency_gate(candidate, allen_relation_map)
    apply_barrier_logic_gate(candidate, protection_logic_context)
    # Candidates failing any gate are moved to ruled_out with the gate reason

# Near-tie and block flagging
for each adjacent pair (c_i, c_{i+1}) sorted by new_composite:
    if (c_i.composite - c_{i+1}.composite) <= review_alternative_gap:
        c_i.review_required = True
        c_{i+1}.review_required = True    # both flagged as near-tie

for each candidate:
    if evidence_posture == "contradicted" OR temporal_posture == "contradicted":
        candidate.review_required = True
        candidate.primary_eligibility = "blocked"
        candidate.primary_block_reasons = [...]

```

causality_candidates (v2) = re-ranked by new_composite

Near-tie and block flags (review_required , primary_eligibility=blocked) directly affect synthesis in Step 6: the synthesizer includes review_required candidates beyond the normal top-N cutoff and will not select a primary_eligibility=blocked candidate as the primary hypothesis regardless of its composite score.

Outputs

Artifact	Description
causality_candidates (v2)	Re-ranked candidates with updated composite scores, evidence citations, near-tie flags, and block status
causality_candidates_pre_refine	Snapshot of v1 scores; used by rank-inversion attention flag and auto-reentry

What can go wrong

- **No evidence passages retrieved** (Step 5a returned empty): the evidence sub-score remains at its Step 4 initial value; re-scoring still runs but does not meaningfully change the ranking. The run_manifest will flag evidence coverage as absent.
- **Rank inversion after refinement**: a candidate that ranked lower in v1 moves above a higher-ranked v1 candidate after evidence is applied. This is expected when evidence strongly favors a previously low-scoring failure mode. The attention flag in Step 6 will call this out explicitly.
- **All top candidates have review_required=True** : evidence is contradictory or no candidate has a clear score lead. The synthesizer will include them all; the analyst must resolve the near-tie manually before writeback.

Step 5d — Auto-reentry

Purpose If the first full pass (Steps 1–5) produced low-confidence results and the pipeline detects that a broader KG context could improve them, automatically re-run portions of the pipeline with expanded parameters. This avoids requiring the analyst to manually trigger a re-run for common coverage-gap scenarios.

Inputs causality_candidates_pre_refine , causality_candidates (v2), evidence_bundle , kg_context , signal_evidence , tskr_patterns , run_context , and all optional Step 5c inputs.

How it works

The pipeline compares v1 and v2 candidate scores and checks the coverage summary for gaps. If the top-ranked candidate’s composite score is below a confidence threshold and the coverage gap is fixable (e.g., the KG neighborhood can be expanded), a reentry is triggered.

```
if enable_auto_reentry AND confidence_gap_detected(pre_refine, post_refine):

    for attempt in 1 .. auto_reentry_max_attempts:

        # Expand KG neighborhood and re-run Steps 1–5
        kg_context (expanded) = re_query_kg(event, wider_hop_limit)
        signal_evidence, tskr_patterns = re_run_steps_1_to_3(kg_context_expanded)
        causality_candidates, evidence_bundle = re_run_steps_4_to_5(tskr_patterns)
        kg_governance = re_assess_governance(kg_context_expanded)

        if top_candidate_score improved sufficiently: break

    reentry_execution = {triggered: True, attempt_count: attempts,
                        kg_governance: refreshed}

else:
    reentry_execution = {triggered: False, attempt_count: 0}
```

The refreshed kg_governance from reentry is what Step 6 uses for KG governance attention flags — it reflects the post-reentry KG state, not the original Step 1 assessment.

Outputs reentry_execution (persisted); updated kg_context , tskr_patterns , causality_candidates , evidence_bundle , and kg_governance if reentry was triggered.

Parameters and defaults

Parameter	Default	Effect
enable_auto_reentry	True	Enables the reentry loop; typically disabled in test/fixture runs for determinism
auto_reentry_max_attempts	1	Maximum number of reentry iterations

What can go wrong

- **Reentry triggered but does not improve scores:** the expanded KG context did not add useful failure modes. The final candidates are from the last reentry iteration; the `reentry_execution` record documents what was attempted.
- **Reentry disabled:** `reentry_execution.attempt_count` is 0 and `triggered` is False. Coverage gaps detected in Step 5c will appear in the manifest as data-limited flags but will not trigger a second pass.

5.12 — Step 6: Conclusion and finalization

Step 6 assembles the final deliverables. It runs eight sub-steps in sequence: (6a) optional Ishikawa 6M bucketing of evidence themes; (6b) barrier analysis aggregating safety function state across all candidates; (6c) final assembly of the similar event list for the synthesizer; (6d) RCA card synthesis from all prior artifacts; (6e) attention flags applied to the card; (6f) output schema validation; (6g) optional Chroma archive; (6h) scope expansion signal detection; and (6i) run manifest finalization, workflow dispatch, and return. The `rca_card` and `run_manifest` produced by this step are the two primary deliverables of the pipeline.

Step 6a — Ishikawa matrix (optional)

Purpose Provide a structured 6M cause-category view of the evidence and candidates, supplementing the failure-mode ranking with a causal-category perspective. This step is optional: if `enable_ishikawa` is False, no matrix is produced and an attention flag is set on the RCA card.

Inputs `event` , `telemetry_summary` , `kg_context` , `tskr_patterns` , `causality_candidates` (v2), `evidence_bundle` , `operational_context` , `pm_compliance` , `run_context` .

How it works The Ishikawa evaluator maps evidence passages and candidate metadata onto six cause categories: personnel / human factors, equipment / hardware, procedure / process, material / parts, measurement / instrumentation, and environment. It produces a matrix of which categories have supporting evidence and which are under-represented.

```
if enable_ishikawa:

    for each evidence_passage and candidate:
        bucket = classify_6M(passage.text, candidate.fm_category)
        ishikawa_matrix[bucket].append(passage_or_candidate)

    # flag under-represented categories
    for each 6M bucket with no entries:
        ishikawa_matrix.gaps.append(bucket)

else:
    ishikawa_matrix = None    # attention flag set at Step 6e
```

Outputs `ishikawa_matrix` (persisted if produced) or `None` .

What can go wrong If `enable_ishikawa` is True but no evaluator is configured, the pipeline raises immediately. If disabled, an Ishikawa-skip attention flag appears on the RCA card noting that the 6M view was not assessed.

Step 6b — Barrier analysis

Purpose Produce a run-level summary of safety function and barrier status — which barriers were intact, which were degraded or failed, and how that relates to the top candidates. This is a run-level aggregation, distinct from the per-candidate `barrier_held` hard gate applied in Step 4.

Inputs `event` , `kg_context` , `causality_candidates` (v2), `evidence_bundle` , `ishikawa_matrix` (optional).

How it works

```
for each safety_function in kg_context.safety_functions:

    # Determine barrier status from protection logic and candidate gate results
    if protection_logic_context has explicit actuation record:
        status = protection_logic_context[safety_function].status
    elif any candidate has hard_gate.barrier_held == True for this function:
        status = "intact"
    else:
        status = "unknown"

    barrier_analysis[safety_function] = {
        status: status, # intact | degraded | failed | unknown
        associated_candidates: candidates linked to this safety function,
        evidence_references: relevant evidence passages
    }

barrier_analysis.degraded_barrier_count = count where status in ("degraded", "failed")
```

Outputs `barrier_analysis` (persisted). A compact summary is injected into the `rca_card` at Step 6d.

Step 6c — Similar event list assembly

Purpose Assemble the final `similar_event_list` and pass it to the RCA card synthesizer. The scoring and querying logic was described in Step 2d (§5.4); this sub-step is the execution point where that list is finalized in pipeline order, immediately before synthesis.

Inputs `event` , `kg_context` , `causality_candidates` (v2). `SimilarEventAdapter` if configured.

How it works The top active candidates drive the component and failure mode IDs used to filter the plant past-event pool and query fleet/industry adapters (same logic as Step 2d). The assembled list is passed directly to the synthesizer in Step 6d and embedded in the run manifest.

Outputs `similar_event_list` — passed to synthesizer; embedded in `run_manifest.artifacts` .

Step 6d — RCA card synthesis

Purpose Assemble the primary analyst-facing deliverable — the `rca_card` — from all artifacts produced in Steps 0–6c. The card presents the ranked failure mode hypotheses, evidence citations, similar event findings, barrier status, and data quality notes at three causal depths: proximate cause, contributing causes, and root cause.

Inputs `causality_candidates` (v2), `evidence_bundle` , `tskr_patterns` , `kg_context` , `telemetry_summary` , `ishikawa_matrix` , `similar_event_list` , `cmms_context` , `operational_context` , `pm_compliance` , `run_context` .

Note on `barrier_analysis` : this is **not** a `synthesize()` input. The barrier summary is injected into the card by the orchestrator **after** synthesis via `rca_card["barrier_analysis"] = _barrier_summary_for_card(barrier_analysis)` .

How it works The synthesizer operates in one of two modes:

- **LLM-assisted** (when an LLM client is configured): `generate_json()` is called with a large structured prompt containing candidates, evidence slices, TSKR patterns, and other inputs. The LLM returns a **full `rca_card` JSON** — including primary hypothesis, contributing causes, evidence citations, recommended actions, and narrative text. Deterministic postprocessing rules then run on top of the LLM output: safety significance, metamodel phase-2, recommended action depth mapping, CCF summary injection, and human performance assessment injection. If the LLM returns a `primary_hypothesis.candidate_id` that is not in the active candidate set, the output is discarded (hallucination guard) and the fallback path is used instead.
- **Deterministic fallback** (`DummyLLMClient` or after hallucination discard): `_fallback_card()` fills the card from the ranked candidate list using rule-based templates. This path is deterministic and is the recommended baseline for regression testing.

similar_event_list and the LLM prompt: `similar_event_list` is accepted by `synthesize()` but is **not included in the LLM prompt** — it is used only by the deterministic fallback path. When the LLM path succeeds, the similar event data reaches the card via the orchestrator's post-synthesis injection steps, not through the LLM itself.

```
# LLM path (when LLM client configured)
prompt = build_prompt(candidates, evidence_slices, tskr_patterns,
                      operational_context, pm_compliance, ...)
# Note: similar_event_list is NOT in the prompt

llm_output = llm_client.generate_json(model, prompt, temperature=0.1)

if llm_output.primary_hypothesis.candidate_id NOT in active_candidate_ids:
    # hallucination guard: discard LLM output, use fallback
    rca_card = _fallback_card(candidates, evidence_bundle, similar_event_list, ...)
else:
    rca_card = llm_output

# Deterministic postprocessing always runs (both LLM and fallback paths):
rca_card = apply_safety_significance_postprocessing(rca_card)
rca_card = apply_metamodel_phase2_postprocessing(rca_card)
rca_card = enforce_recommended_action_depth_mapping(rca_card)
rca_card["ccf_summary"] = build_ccf_summary(causality_candidates)
rca_card["human_performance_assessment"] = build_hp_assessment(ishikawa_matrix)

# Orchestrator injects barrier summary after synthesize() returns:
rca_card["barrier_analysis"] = barrier_summary_for_card(barrier_analysis)
```

Outputs `rca_card` (persisted; primary return artifact).

Parameters and defaults

Parameter	Default	Effect
LLM client	None (<code>DummyLLMClient</code> raises, fallback used)	When a real LLM client is configured, the full card JSON is LLM-generated then deterministically postprocessed
<code>allow_fallback_template_fill</code>	True	If <code>True</code> , LLM failure or hallucination guard triggers the deterministic fallback rather than aborting

Step 6e — Attention flags

Purpose Apply a set of rule-based checks to the synthesized `rca_card` and annotate it with flags that draw the analyst’s attention to specific quality or data concerns. Flags do not change the ranking; they add a structured alert to the card.

How it works Seven flag checks are applied in sequence after synthesis:

```

# 1. Rank inversion: post-refine ranking differs from pre-refine for top candidates
if top_candidate changed between causality_candidates_pre_refine and v2:
    add_flag("rank_inversion", {pre_rank, post_rank, candidate_id})

# 2. KG governance: KG was amber or red state
if kg_governance.status in ("amber", "red"):
    add_flag("kg_governance", kg_governance.issues)

# 3. Recurrence match quality: top candidate has unresolved recurring failures
if top_candidate.recurrence_count > 0 AND top_candidate.unresolved_recurrence_count > 0:
    add_flag("recurrence_match_quality", {recurrence_count, trend, unresolved_count})

# 4. Ishikawa skipped: matrix was not produced
if ishakawa_matrix is None:
    add_flag("ishikawa_skipped")

# 5. Signal evidence gaps: signal tags with no KG component link
if signal_evidence has unresolved tags:
    add_flag("signal_evidence_gap", {unresolved_tag_count})

# 6. Out-of-boundary signals: anomaly signals outside current KG scope
if any signal.component_id NOT in kg_context.components:
    add_flag("out_of_boundary", {out_of_scope_signals})

# 7. Metamodel coverage: causal categories A-L under-represented in candidates
if coverage(categories_A_to_L) < threshold:
    add_flag("metamodel_coverage", {under_represented_categories})

```

Outputs `rca_card` updated in-place with `attention_flags` list; persisted.

Step 6f — Output validation

Purpose Validate all output artifacts against their JSON schemas — the same schema check applied to inputs at Step 0, now applied to outputs. This catches schema drift when the pipeline emits fields the schema does not yet know about, which is the primary signal that a schema update is needed after a code change.

Inputs All output artifacts: `rca_card`, `causality_candidates`, `evidence_bundle`, `barrier_analysis`, `ishikawa_matrix`, `tskr_patterns`, `kg_context`, `signal_evidence`, `input_validation`.

How it works

```

output_validation = validate_all_outputs(
    rca_card, causality_candidates, evidence_bundle,
    barrier_analysis, ishakawa_matrix, ...
)

for each artifact:
    if schema_check fails AND artifact is required:
        if stop_on_validation_error: abort run
        else: record in output_validation.failures
    if schema_check fails AND artifact is optional:
        record in output_validation.optional_failures

```

Outputs `output_validation` (per-artifact pass/fail results).

Step 6g — Chroma archive

Purpose Write the run's evidence passages and embeddings back to a long-lived Chroma collection so that future runs can retrieve them as operating experience. This step is deployment-specific and is typically disabled in test environments.

How it works

```
if enable_chroma_archive_stage:

    archive_result = write_to_chroma(run_id, evidence_bundle, rca_card)

    if archive_result.failed AND hard_fail_on_chroma_archive_error:
        mark run_status = "aborted"
        raise error    # run aborts AFTER rca_card is already persisted

else:
    archive_result = {enabled: False}    # no-op
```

Outputs chroma_archive result dict (status, written count, or error).

Parameters and defaults

Parameter	Default	Effect
enable_chroma_archive_stage	False (test); True (production)	Enables the archive write
hard_fail_on_chroma_archive_error	True	If True , archive failure aborts the run after the RCA card is already saved

What can go wrong An archive failure under strict policy aborts the run and marks run_status = "aborted" — but the rca_card is already persisted and accessible. Set hard_fail_on_chroma_archive_error: False to treat archive failures as non-fatal warnings.

Step 6h — Scope expansion signals

Purpose After the RCA card is finalized, scan the Allen relation map, signal evidence, and TSKR patterns for signals from components that are outside the current investigation scope but show precursor timing. Any such signals are recorded as *suggested scope expansions* that the analyst can accept before the next run. This step does not change the current run’s candidates — it only seeds the next run’s scope.

Inputs run_context , allen_relation_map , signal_evidence , tskr_patterns .

How it works

```
expansion_signals = []

for each signal in (allen_relation_map + tskr_patterns):
    if signal.component_id NOT in current_scope.components AND
       signal.allen_relation == "PRECEDES":    # potential cause outside scope
        expansion_signals.append({
            component_id: signal.component_id,
            reason:       "out-of-scope precursor signal",
            evidence:     signal.confidence / allen_base_score
        })

if expansion_signals:
    run_context.scope_management.expansion_suggestions = expansion_signals
    # Analyst reviews suggestions and accepts/rejects via the scope revision API
    # Accepted suggestions create scope version N+1 for the next run (see §2.5.3)
```

Outputs Updated run_context (re-persisted if expansion signals were found). Expansion suggestions are visible in run_context.scope_management.expansion_suggestions .

Step 6i — Run manifest finalization and return

Purpose Assemble the complete audit record for the run (`run_manifest`), execute workflow dispatch if configured, mark the run complete, and return all artifacts to the caller.

Inputs All artifacts produced during the run; `allen_relation_map` ; `similar_event_list` .

How it works

```
# Build signal_lessons_learned (Step 3.5 output – computed here at finalization)
signal_lessons_learned = classify_patterns(tskr_patterns) # matched vs novel per FM

# Assemble run_manifest
run_manifest = {
    artifacts: {
        data_coverage_summary: coverage of all input dimensions (kg, tskr, evidence, ...),
        sensitivity_table:      which inputs, if changed, would most affect the ranking,
        scope_state:           active_scope_version, expansion_suggestions, filter history,
        signal_lessons_learned: matched/novel classification per failure mode (Step 3.5),
        similar_event_summary: summary of similar_event_list findings
    },
    review_hooks: {
        writeback_targets:      systems to notify (CMMS, DMS, workflow queue),
        human_review_required:  True if any attention flag is present,
        decision_status:        rca_card.decision_status,
        ...                     (28 fields total)
    },
    pipeline_config:           configuration snapshot for reproducibility
}

# Workflow dispatch (runs after manifest is assembled, merges results back)
if workflow_dispatch configured:
    dispatch_result = execute_workflow_dispatch(run_manifest)
    run_manifest.review_hooks += dispatch_result
    run_manifest.artifacts    += dispatch_result.artifacts

# Persist and mark complete
save(run_manifest)
save(run_status, run_complete=True)

return {
    run_context, pm_compliance, kg_context, signal_evidence, tskr_patterns,
    causality_candidates, causality_candidates_pre_refine, evidence_bundle,
    ishikawa_matrix, barrier_analysis, reentry_execution, cmms_context,
    rca_card, input_validation, output_validation, run_manifest
}
```

Outputs `run_manifest` (persisted); complete return dict with all artifacts listed above. `run_id` is the correlation key across all artifacts, logs, and the artifact store.

What can go wrong

- **Step 3.5 (`signal_lessons_learned`) empty:** no TSKR patterns or all patterns have `confidence = 0` . The manifest will note TSKR data as absent.
- **Workflow dispatch fails:** logged and skipped; `run_manifest` is saved without dispatch metadata. The run is still marked complete.
- `run_status.run_complete` **remains False** : the run did not reach this step — check logs for which earlier step failed.

6. Output Artifacts (Reader-Oriented)

Artifact	Manager use	System engineer use	Data analyst use
rca_card	One-screen primary story, gaps, decision status	Trace candidate IDs, categories, attention flags	N/A (unless curating features)
run_manifest	Review hooks, run health, scope and sensitivity	Stage health, artifacts sub-objects, validation links	Data coverage and missing source list
causality_candidates / _pre_refine	N/A (detail)	Scores, gates, chain_position	Provenance to raw inputs
evidence_bundle	N/A (unless cited in exec summary)	Snippet / doc alignment	Chroma query tuning
tskr_patterns	“Something preceded something” in plain language	Pattern IDs, Allen support	TSKR data feed QA
ishikawa_matrix	Programmatic / procedure gaps	Category rows and keywords	N/A
barrier_analysis	“Which safety function failed” at a high level	PLC consult flags, degraded_barrier_count	N/A
input_validation / output_validation	“Safe to use for review”	Schema failures	Schema drift after upgrades
reentry_execution	N/A (ops / tuning)	Whether reentry ran, kg / candidate mutations and hook output	N/A
pm_compliance / cmms_context	PM and WO narrative when present	Build provenance (auto vs provided), CMMS-augmented KG	Source alignment for PM rows

Reading data_coverage_summary and sensitivity_table (both under run_manifest.artifacts): these are the primary analyst view of "what we had" vs "what could have changed the answer."

Status	Meaning
not_assessed	Input was not passed to run() at all — distinct from passing an empty value
missing	Passed but absent or could not be retrieved
partial	Present but incomplete or quality-flagged

If SOE is present and PLC is absent, SOE/PLC pairing may be violated — see the TSKR chain bonus (TC-2). Cross-reference: **TC-7** (degraded SOE/alarm path), **TC-2** (PLC pairing).

Note: similar_event_list , signal_lessons_learned , and workflow_dispatch details live under run_manifest.artifacts / review_hooks ; run_id is the correlation id in the artifact store and run_status (see §5.18).

7. Human and Organizational Contributors (G / I / L)

- G** (execution error) and **I** (procedure / config baseline wrong) are **distinguishable** in principle; the product surfaces both when evidence supports them (**TC-6**).
- Training records** establish **qualification and recency** — they can **rule in** a human-factors *investigation* but do not by themselves establish root cause.
- L** (organizational) requires documentary and often **OE** support — use SimilarEventAdapter and fleet data when available.

Show-and-tell cross-reference: TC-6.

8. Appropriate Use and Limitations

- The pipeline does **not** perform physical inspection, offline lab analysis, or NRC reporting determinations. It is a decision-support tool; final determinations remain with the responsible engineer or review board.
- **Scores and gates are model output.** Challenging a ranking or gate decision with new evidence is the normal and expected process — the manifest is explicitly designed to surface gaps so that reviewers can do this.
- **Reproducibility:** Given the same code version and the same inputs, all rule-based and retrieval stages produce identical outputs. LLM-synthesized narrative (the `rca_card` text and similar fields) may vary across runs when a live LLM is enabled; the Dummy LLM path produces stable text and is the recommended baseline for regression testing. See **Appendix B** for a precise boundary between LLM-generated and rule-generated content.
- **Known limitations** in pattern recognition and similar-event search are documented in **Appendix A** and **Appendix C** respectively — consult these before drawing conclusions from low-confidence or novel-pattern runs.

Appendix A — Pattern recognition: limitations and improvement directions [WIP]

Section §2.4 describes what pattern recognition means in this pipeline, how each mechanism works, and what it produces. This appendix covers what §2.4 does not: where the pattern logic can misfire, and where the product can be deepened.

Debugging note: FM-anchored patterns (TSKR), event-wide signal order (Allen), hypothesis ranking (causality engine), and organizational recurrence (similar events) are *different* objects backed by different artifacts. Do not conflate them when tracing a surprising candidate rank or card narrative — each has its own inputs-outputs path in §2.4 and §5.

A.1 Where pattern logic can misfire (limitations)

- **Data boundedness:** TSKR and Allen are only as good as the **time windows and logs** you pass in. Missing SOE, sparse alarms, or a thin `telemetry_summary` yield **low temporal support** without labeling the run “invalid” — the manifest and attention flags are how gaps surface.
- **KG coverage:** If the hypothesis space omits a relevant FM or the graph’s **past_events** are incomplete, *recurrence* and *signature-FM* matches skew toward the **known** set, not the **true** set.
- **Scorer is engineered, not learned:** `CausalityEngineConfigV32` **weights and caps** (e.g. CCF/OP deltas) are hand-tuned for interpretability. They are **not** a trained neural end-to-end model; changing plant priorities may require **config and governance** review, not retraining a single model.
- **No substitute for physics or on-site work:** Distinguishing *instrumentation* vs *real* plant response, or a one-off from a systematic failure, still needs human judgment and often **work management / inspection** data outside this pipeline.
- **TSKR absent path:** If `tskr_temporal_scorer` is `None`, patterns are **synthetic / empty** — downstream still runs, but metamodel *temporal pattern* questions are under-supported.

A.2 Room for improvement (product and engineering, not promises)

- **Richer TSKR inputs:** Tighter integration with the plant historian and alarm **pairing** quality would raise confidence in relation labels without changing the *architecture* of §5.4.
- **Calibration and feedback:** Closed-loop use of analyst **accept / reject** on primary hypotheses could inform **weight** or **gate** policy over time, while keeping **deterministic replay** a requirement.
- **OE depth: Appendix C** — stronger fleet/industry retrieval and **schema** for `similar_event_list` (see `step2d_similar_event_plan_april_25.md`) would strengthen the “have we seen this before?” path.
- **Scope and governance:** Tighter **KG governance** and **scope** UX reduce false confidence when the graph is stale or the investigation boundary is still fluid (**TC-7** narrative).

Appendix B — Large language models: roles today and possible futures

B.1 What uses an LLM *today* (in this repo)

- **RCA card structured synthesis** — `RuleValidatedRCASynthesizerV31` calls `llm_client.generate_json(...)` to turn a **large, tabular** prompt (candidates, evidence slices, TSKR, Ishikawa, etc.) into **rca_card JSON** (`synthesis/rca_synthesizer_v31.py`). If generation fails, validation fails, or the model **invents** a `primary_hypothesis.candidate_id` that is not in the input set, the synthesizer **discards** the LLM output and uses `_fallback_card` (deterministic template fill) when `allow_fallback_template_fill` is `True`.

- **Optional similar-event (fleet / industry) tier** — A deployment may inject `LLM0EAdapter` (`adapters/llm_oe_adapter.py`) as a `SimilarEventAdapter` implementation. That is an **HTTP API to an LLM- or RAG-style service**, not a call inside the causality engine. Plant-tier matching (Appendix C) is **in-memory, non-LLM**.

B.2 What does *not* use the narrative LLM by default

- **Candidate ranking and hard gates** — `RuleBasedCausalityEngineV32` (rule + score, §5.6).
- **Vector retrieval (Chroma)** — Uses **embeddings** and retrieval config (§5.8); the embedding model is a **separate** concern from the **synthesis** LLM. Do not assume one vendor or one model for both.
- **TSKR temporal scoring** — External scorer or synthetic absence (§5.4).
- **Ishikawa, barrier aggregation, input validation, manifest math** — Deterministic in Python.

B.3 Guardrails and determinism (SE-relevant)

- **DummyLLMClient** — `generate_json` **raises** on purpose so development runs use **fallback** `synthesis` (`orchestrators/llm_clients.py`).
- **Hallucination guard** — If the model picks a **primary** `candidate_id` that was **never** an input candidate, the LLM `rca_card` is dropped (see `RuleValidatedRCASynthesizerV31.synthesize`).
- **Post-processing is deterministic** — After either LLM or fallback path, the synthesizer applies **safety / metamodel** postprocessing, **injects** `ccf_summary` and may inject `human_performance_assessment` in code, so the **safety story** is not fully delegated to the LLM.
- **Narrative variance** — With a real LLM (e.g. `Ollama` via `ollamaLLMClient`), **wording** and **ordering** in executive text may differ run-to-run at fixed inputs; **IDs** and **score plumbing** are still constrained as above.

B.4 Possible futures (design space, not committed roadmap)

- **Assisted red-team / second pass** — A model could critique **gaps** or **contradictions** only after **deterministic** candidates exist, to avoid conflating “safety logic” with “chat.”
- **Stronger RAG** — Tighter coupling between **Chroma** chunks and **citations** in the card, with explicit *no-source* rules for sensitive claims.
- **Governed generation** — Fixed JSON schema, **field-level** allowlists, and temperature **0** for ID-bearing fields; multi-stage generate-validate.
- **Fleet OE** — LLM-backed or hybrid retrieval (Appendix C) for industry databases with **provenance** and **tier discount** as first-class fields.

Appendix C — Similar event search in the RCA workflow

This is the place **SEs** will look to answer: *How do we “search” for similar events, and why might we miss a known industry case?*

C.1 When it runs and what it produces

- **Stage: After** `barrier_analysis` and **before** `rca_synthesizer.synthesize` (§5.12, sub-steps 6b → 6c → synthesis). The list is passed as `similar_event_list` so the **card** and **narrative** can refer to recurrence and `unresolved_gaps` .
- **Artifact:** A single `similar_event_list` **dict** with `status` , `query_terms` , `summary` (counts, `degraded_tiers` , `any_plant_match`), `events[]` , and `provenance` . The **return dict** of `run()` does **not** top-level a separate key for it — the **manifest** and card consumption carry it; do not assume a parallel return field in strict JSON consumers.

C.2 Query terms (audit trail)

- Built from the **top N causality** candidates (default `step2d_query_top_n_candidates = 3` in `OrchestratorConfig.extra` — see `_build_similar_event_list`).
- **Extracts:** `component_id` , `failure_mode_id` (or `canonical_tuple.failure_mode` fallback), and event-level `event_type` and `actuation_type` , plus `asset_id` — captured under `query_terms` for **traceability** (not for hidden ML).

C.3 Plant tier (always on)

- **Source:** `kg_context.past_events` (temporal metadata — `in_precursor_window` and related fields — is injected by `_enrich_past_events_temporal_metadata` at the end of Step 1, after CMMS augmentation).
- **Scoring** (`_query_plant_past_events`): for each past event, **partial credit** is summed from:

Dimension	Condition	Weight
Component match	<code>matched_component_ids</code> non-empty for that past event	0.40
FM match	intersection of top candidates' FM IDs and <code>matched_failure_mode_ids</code>	0.25
Event type	string equality of <code>event_type</code> to current	0.15
Actuation type	string equality of <code>actuation_type</code> to current (if current set)	0.10
Precursor window	<code>in_precursor_window</code> <code>True</code>	0.10

- The raw sum is **capped** at 1.0, then multiplied by the **plant** tier factor **1.0** from `TIER_CONFIDENCE_MULTIPLIERS` (`adapters/similar_event_adapter.py`). Results are **sorted** by `confidence_weight` and trimmed to `step2d_plant_top_n` (default **5**). **This is a structured heuristic, not a semantic text search of free-text LERs** unless a higher layer puts text into `past_events` with the right fields.

C.4 Fleet and industry tiers (optional)

- **Precondition:** `orchestrator.set_similar_event_adapter(adapter)` ; default is **no adapter** → fleet and industry are **skipped**, `status` = **partial** , and provenance notes *no adapter injected*.
- **Contract:** `SimilarEventAdapter.query(level="fleet"|"industry", ...)` returns a list of **dict** rows; the orchestrator must **not** get exceptions from well-behaved adapters (errors → empty list and `degraded_tiers` metadata). Scores are adjusted by **tier multipliers** — e.g. **fleet 0.80**, **industry 0.60** — to avoid over-crediting external digests.
- **Reference implementation:** `LLMOEAdapter` — HTTP-based, described in `adapters/llm_oe_adapter.py` and the Step 2d plan. Treat as **deployment-specific wiring**.

C.5 Status, degradation, and SE interpretation

- **complete** only when an adapter was present and **neither** tier is marked degraded. Otherwise **partial** . Either way, **zero plant matches** with no fleet/industry results means “**no similar event surfaced**” in **product terms**, not “none exist in the world.”
- **SE action:** If recurrence is business-critical, validate **past_events quality** in the KG, **increase** `step2d_plant_top_n` / candidate query width only with awareness of **noise**, and invest in **adapter** + source registry — not a larger LLM in the main orchestrator by itself.

C.6 Further reading in-repo

- `diagrams/april_25/step2d_similar_event_plan_april_25.md` — product and schema evolution for Step 2d.
- **Show-and-tell: TC-7** (`any_plant_match`); **TC-3** where `past_events` is populated in the KG fixture.

Appendix D — Show-and-tell test case index

Case	File	Emphasis
TC-1	<code>tests/test_case_1/</code>	Minimal plumbing, <code>build_dev_orchestrator</code> path optional
TC-2	<code>tests/test_case_2/</code>	SOE, alarms, PLC pairing, vacuum loss
TC-3	<code>tests/test_case_3/</code>	Condenser / U2 scenario, configuration + environmental
TC-4	<code>tests/test_case_4/</code>	Reactor trip, timeline gate , RPS / PLC
TC-5	<code>tests/test_case_5/</code>	ECCS CCF, vendor supply chain
TC-6	<code>tests/test_case_6/</code>	Human performance, training
TC-7	<code>tests/test_case_7/</code>	Scope expansion , TSKR, degraded SOE/alarm (absent)

Document history

Version	Date	Notes
0.1	2026-04-25	Initial full draft; AP-913 deferred; code anchors for Neo4j/Chroma/test paths.
0.2	2026-04-25	Section 5 expanded to 18 execution-ordered substeps (5.1–5.18); PM build order vs KG clarified.
0.3	2026-04-25	§2.3 Pattern recognition (metamodel table + <code>signal_lessons_learned</code>); backlog table row; cross-refs in §1, §2.2, 5.3–5.5, 5.8, 5.13, 5.18 ; former §2.3–2.4 renumbered to 2.4–2.5 .
0.4	2026-04-25	§5 code-order verification blurb; 5.10 reentry / <code>kg_governance</code> ; 5.18 workflow-dispatch <i>after</i> finalize; Appendices A–D (pattern recognition, LLMs, similar events, test index renumbered from former Appendix A).
0.5	2026-04-25	<code>check_doc_code_sync.py</code> first added under <code>scripts/</code> ; §0 with run instructions; <code>@code</code> links verified.
0.6	2026-04-25	Final code alignment with <code>RCAReasoningOrchestrator.run</code> (§1 return bundle, §2.2 step order incl. Chroma / expansion / manifest / return, §2.3 v32 past event hook, §8 reentry / PM/CMMS / manifest note); new §1.1 central SE element table.
0.7	2026-04-25	<code>check_doc_code_sync.py</code> moved to <code>diagrams/april_25/</code> ; §0 <code>@code</code> and run command updated.